

# ICSI 403

# DESIGN AND ANALYSIS OF ALGORITHMS

## Chapter 6 - Heaps, Heapsort, and Priority Queues

# Heapsort - Introduction

- ⦿ Introduces a new data structure (the heap)
- ⦿ Combines best of Merge Sort & Insertion Sort
- ⦿  $O(n \lg n)$  (guaranteed), like Merge Sort
  - Better than Insertion Sort's  $O(n^2)$
- ⦿ In-place sort like Insertion Sort
  - Requires only a constant amount of storage beyond the array to sort, unlike Merge Sort, which requires  $O(n)$  additional storage space, or Quicksort (which requires stack space)
- ⦿ Not as fast (in practice) as Quicksort

# Introduction

## ◎ “A Heap”

- Originally derived from Heapsort, but later used to refer to the garbage-collected data storage pools in LISP, C, and Java
  - Some other languages use garbage-collected heaps, too
- Don't get the two confused – Heapsort has nothing to do with garbage-collected data storage

## ◎ There are “Min Heaps” and “Max Heaps”

- Details to follow

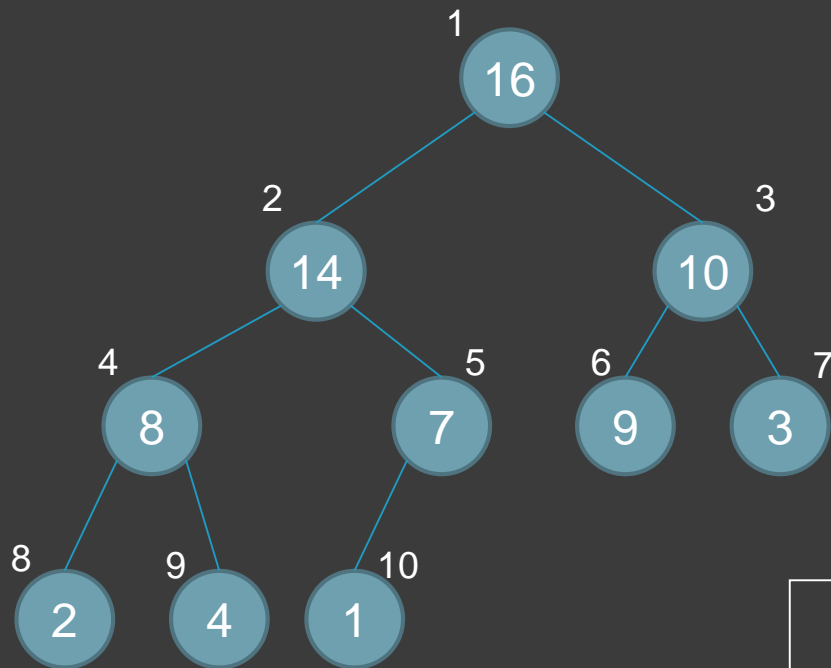
# A Heap

- ⦿ A nearly-full binary tree
- ⦿ NOT a Binary Search Tree (BST)
  - We don't go down through the tree comparing keys to find information.
- ⦿ The tree is completely full, except for the right part of the lowest level (the left part of the lowest level IS completely full, up to a point)
- ⦿ A simple array can be used to store a heap – we don't need tree nodes and pointers!

# Storing a Heap in an Array

- ⦿ Suppose we store a heap in array  $A$
- ⦿ We have two pieces of information to keep track of regarding  $A$ 
  - $A$ 's length (how large is  $A$ ) –  $A.length$
  - $A$ 's heapsize (how many items in  $A$  are part of the heap) –  $A.heapsize$ 
    - Not everything in  $A$  is required to be part of the heap (part of what's in  $A$  can be outside of the heap – we'll come back to this later)
    - There are  $A.length - A.heapsize$  elements of  $A$  that are not part of the heap

# An Example Heap



1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1

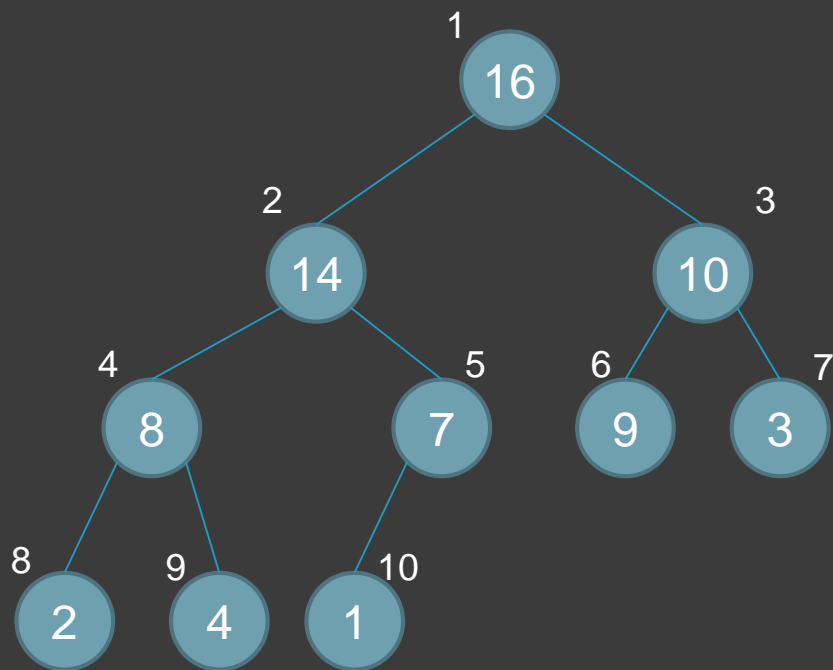
If we number the nodes  
top-to-bottom, left-to-right, ...

We get the subscripts for  
storing the heap in an array

The root is stored in  $A[1]$ . Heaps are not  
typically stored as zero-based arrays

For a given node in the tree (i.e., member of array  $A$ ), whose subscript (index) is  $i$   
The subscript (index) of its left child is  $2i$   
The subscript (index) of its right child is  $2i+1$   
The subscript (index) of its parent is  $\lfloor i/2 \rfloor$

# An Example Heap



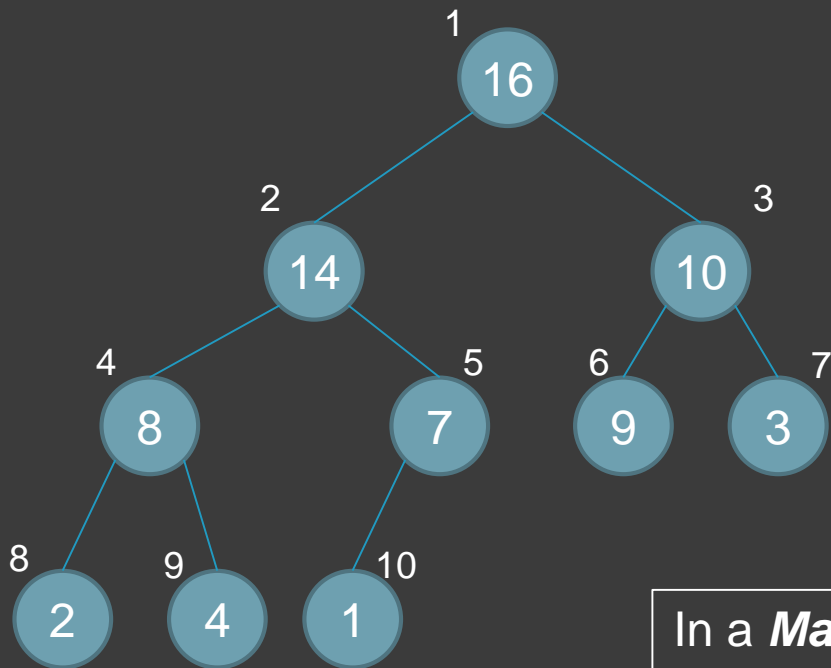
1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1

$2i$ ,  $2i+1$ , and  $\lfloor i/2 \rfloor$  are efficiently calculated using left and right shift operations (and an add)

For a given node in the tree (i.e., member of array  $A$ ), whose subscript (index) is  $i$

- The subscript (index) of its left child is  $2i \rightarrow A[i].\text{LEFT} = A[2i]$
- The subscript (index) of its right child is  $2i+1 \rightarrow A[i].\text{RIGHT} = A[2i+1]$
- The subscript (index) of its parent is  $\lfloor i/2 \rfloor \rightarrow A[i].\text{PARENT} = A[\lfloor i/2 \rfloor]$

# An Example Heap



1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1

For a **Min-Heap**, simply change  $\geq$  to  $\leq$

We said there are two kinds of heaps: Max-heaps and Min-heaps, which satisfy either the Max-Heap or Min-Heap property

In a **Max-Heap**, for every node  $i$  other than the root:  
 $A[i].\text{PARENT} \geq A[i]$   
In other words, the value at a node is, at most, the value of its parent.

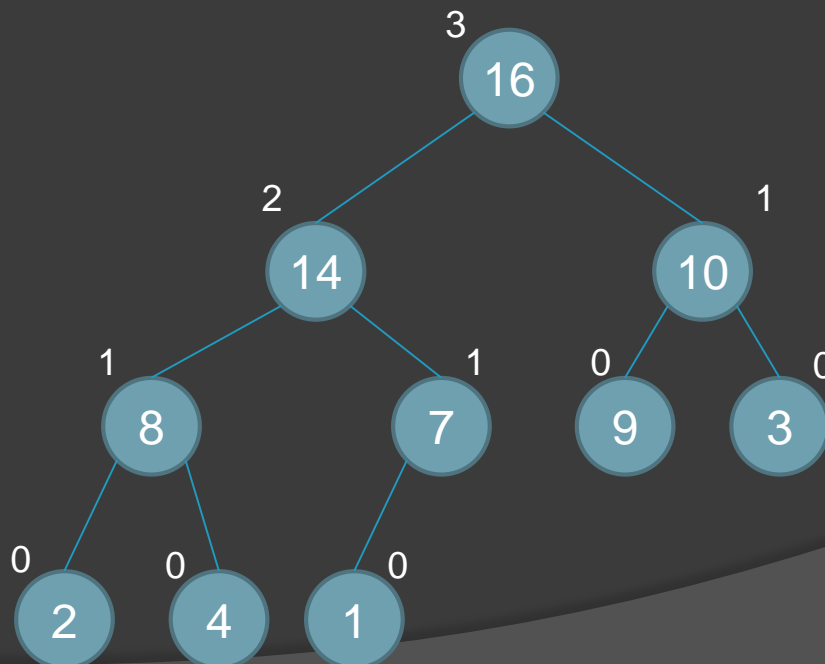
Alternatively, the value at a node is, at least, as large as its child(ren).

For a **Max-Heap**, the largest value is always at the root, or  $A[1]$



# Height of a Node in a Heap

- The height of a node in a heap is the number of edges (child links) on the longest downward path from that node to a leaf
- The height of the entire heap is the height of the root



# Height of a Heap

- Since a heap of  $n$  elements is based on a complete binary tree, the height of a heap is  $\Theta(\lg n)$
- Basic operations on heaps run in time proportional to (at most) the height of the tree, and therefore take  $\Theta(\lg n)$  time.

# Heap Operations on a Max-Heap

- ⦿ MAX-HEAPIFY – runs in  $O(\lg n)$  time – used to maintain the Max-Heap property
- ⦿ Before we run MAX-HEAPIFY,  $A[i]$  may be smaller than its children
- ⦿ Assumes that the left and right subtrees of  $i$  are heaps
- ⦿ MAX-HEAPIFY lets  $A[i]$  “float down” the heap while  $A[i] < \text{its children}$
- ⦿ After MAX-HEAPIFY runs, the left and right subtrees rooted at  $i$  are Max-Heaps

# MAX-HEAPIFY Pseudocode

MAX-HEAPIFY( $A, i$ )

```
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heapsizesize}$  and  $A[l] > A[i]$ 
4       $largest = l$ 
5  else  $largest = i$ 
6  if  $r \leq A.\text{heapsizesize}$  and  $A[r] > A[largest]$ 
7      then  $largest = r$ 
8  if  $largest \neq i$ 
9      exchange  $A[i] \leftrightarrow A[largest]$ 
10  MAX-HEAPIFY( $A, largest$ )
```

Find the  
largest of  
 $A[i]$  and its  
children

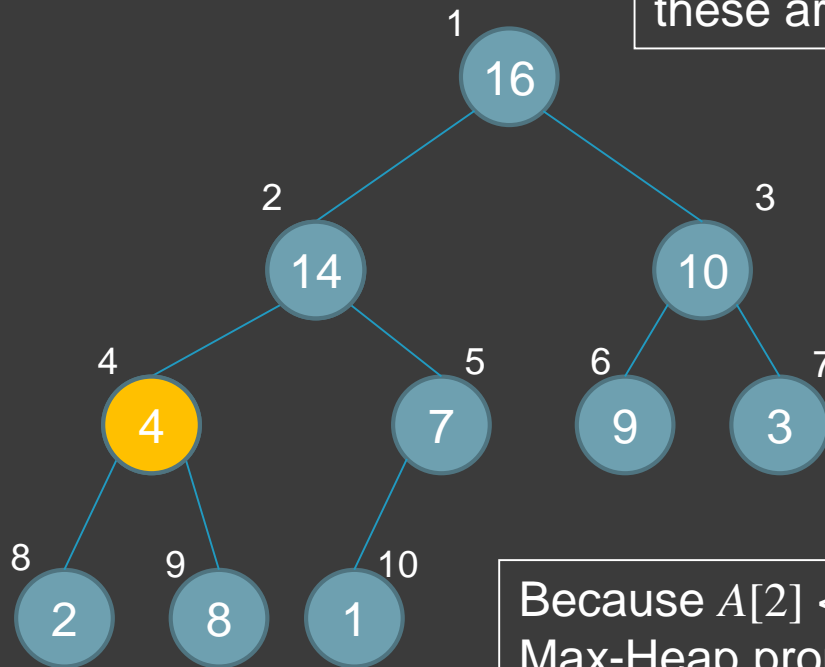
# How MAX-HEAPIFY Works

- ⦿ Compare  $A[i]$ ,  $A[A[i].LEFT]$ , and  $A[A[i].RIGHT]$
- ⦿ If necessary, swap  $A[i]$  with the larger of its children to preserve the Max-heap property
- ⦿ Keep comparing and swapping down the heap, until the subtree rooted at  $i$  is a Max-heap.
  - If we hit a leaf, then the subtree rooted at the leaf is trivially a Max-heap.
- ⦿ Since we only need to go through it once for every level, the running time is  $O(\lg n)$

# Example

## ● MAX-HEAPIFY(A, 2)

Remember, we're really just moving items around in an array; these are not nodes with pointers!



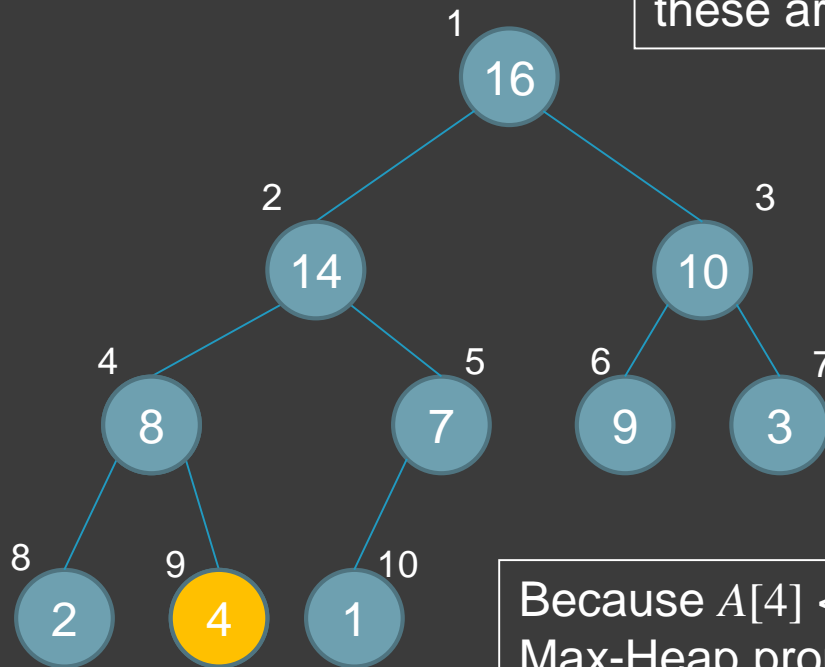
Because  $A[2] <$  one of its children, the Max-Heap property is violated at  $A[2]$

The larger of  $A[2]$ 's children is  $A[4]$  (with a value of 14), so we swap  $A[2]$  and  $A[4]$ , and call MAX-HEAPIFY(A, 4), because we may have just destroyed the Max-Heap property at  $A[4]$ .

# Example

## ⦿ MAX-HEAPIFY(A, 4)

Remember, we're really just moving items around in an array; these are not nodes with pointers!



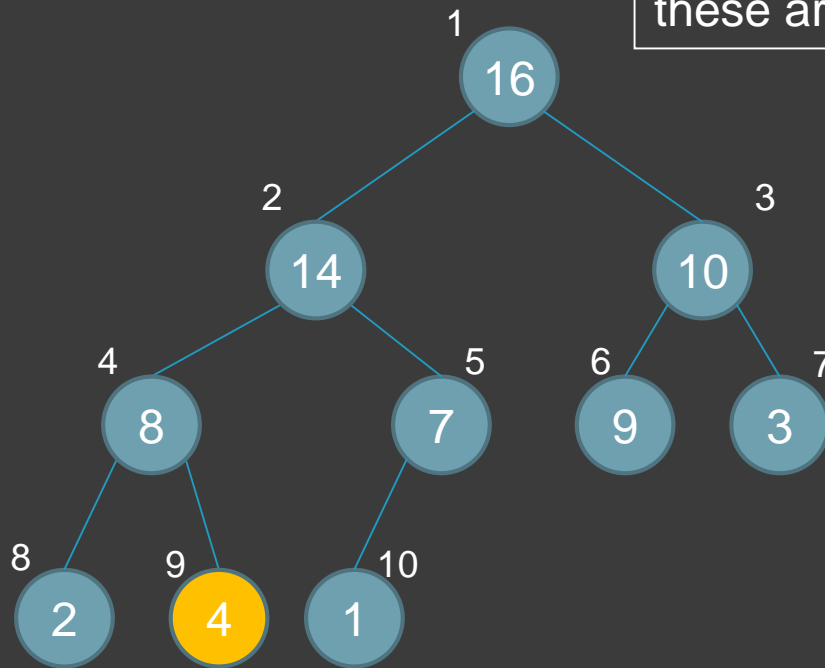
Because  $A[4] <$  one of its children, the Max-Heap property is violated at  $A[4]$

The larger of  $A[4]$ 's children is  $A[9]$  (with a value of 8), so we swap  $A[4]$  and  $A[9]$ , and call MAX-HEAPIFY(A, 9), because we may have just destroyed the Max-Heap property at  $A[9]$ .

# Example

## ⦿ MAX-HEAPIFY(A, 9)

Remember, we're really just moving items around in an array; these are not nodes with pointers!



Because  $A[9]$  is a leaf, it is a trivial Max-heap, and MAX-HEAPIFY ends.  
Because no statements in MAX-HEAPIFY follow the recursive call to MAX-HEAPIFY, the whole process ends.



# Building a Heap

- ⦿ We can use MAX-HEAPIFY to make a heap out of an array from the bottom up
- ⦿ Suppose we have an array  $A[1..n]$ , where  $n = A.LENGTH$
- ⦿ Each leaf is already a heap, so there's no sense running MAX-HEAPIFY on the leaves.
- ⦿ MAX-HEAPIFY can take two subtrees and their parent and make them a Max-Heap
- ⦿ We run MAX-HEAPIFY from all leaves going upward, and then we have a full heap

# Build-Max-Heap Pseudocode

BUILD-MAX-HEAP( $A$ )

- 1  $A.heapsize = A.length$
- 2 **for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1
- 3     do MAX-HEAPIFY( $A, i$ )

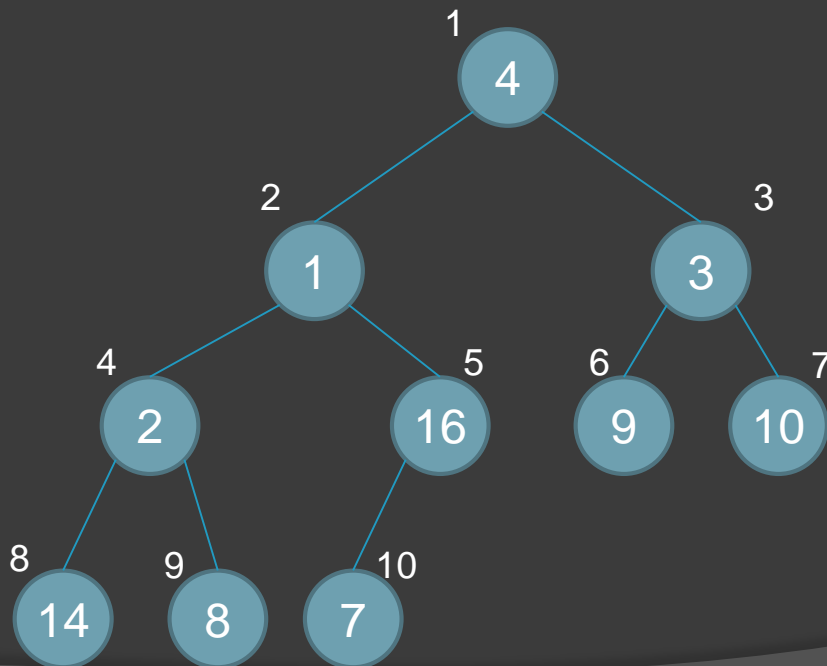
Note:  $\lfloor A.length/2 \rfloor$  is just the left half of the array (less one, if the size of the array is odd)

# Example

- Consider the following 10-element array  $A$

1	2	3	4	5	6	7	8	9	10
4	1	3	2	16	9	10	14	8	7

- If we represent this array as a heap, we have



The BUILD-MAX-HEAP algorithm says:

**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1

**do** MAX-HEAPIFY( $A, i$ )

**for**  $i = \lfloor 10/2 \rfloor$  **downto** 1

Why do we only MAX-HEAPIFY from  $A[5]$  down?

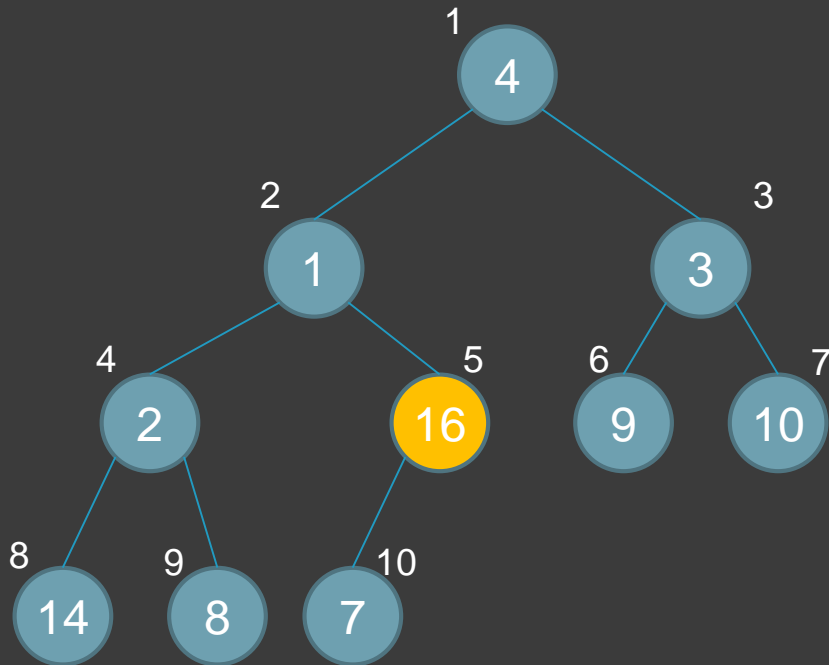
# Example

1	2	3	4	5	6	7	8	9	10
4	1	3	2	16	9	10	14	8	7

**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1  
  **do** MAX-HEAPIFY( $A, i$ )

$A[5]$  (16) is already  $\geq$  its child(ren),  
so there's nothing to do here.

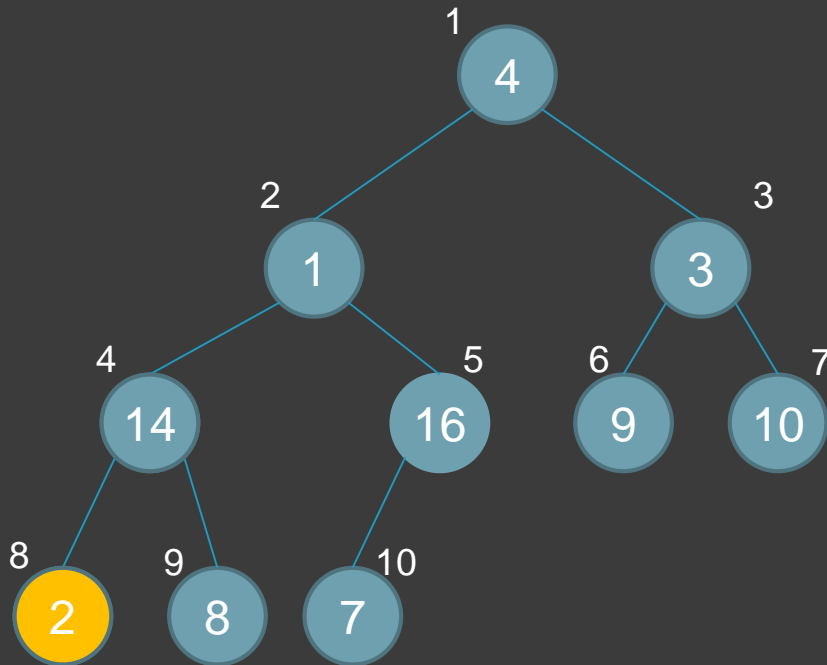
Go on to  $A[4]$



# Example

1	2	3	4	5	6	7	8	9	10
4	1	3	14	16	9	10	2	8	7

**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1  
  **do** MAX-HEAPIFY( $A, i$ )



$A[4]$  (2) is  $<$  its largest child (14), so  
swap  $A[4]$  and  $A[8]$

Then we MAX-HEAPIFY( $A, 8$ )  
 $A[8]$  is a leaf, so we're done here  
Go on to  $A[3]$

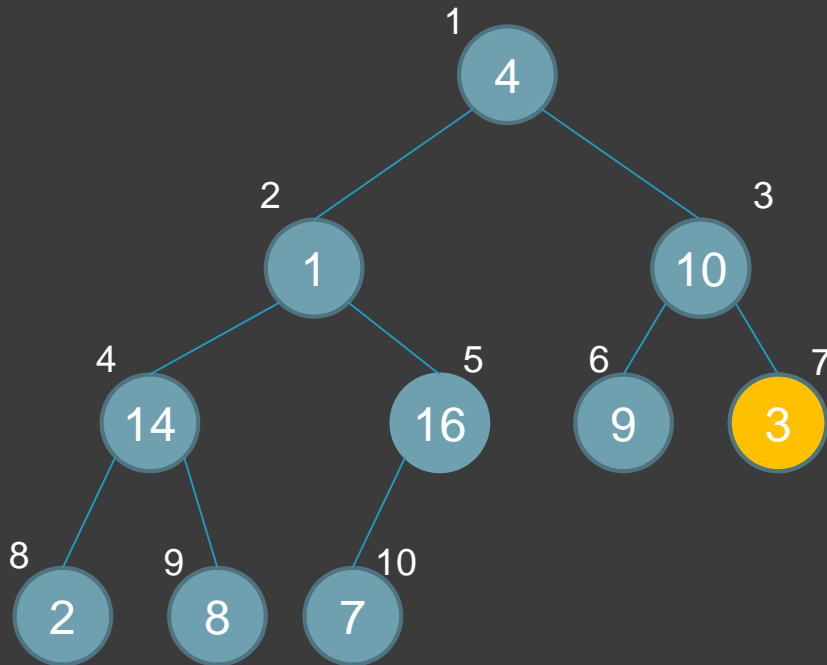
# Example

1	2	3	4	5	6	7	8	9	10
4	1	10	14	16	9	3	2	8	7

**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1  
  **do** MAX-HEAPIFY( $A, i$ )

$A[3]$  (3) is  $<$  its largest child (10), so  
swap  $A[3]$  and  $A[7]$

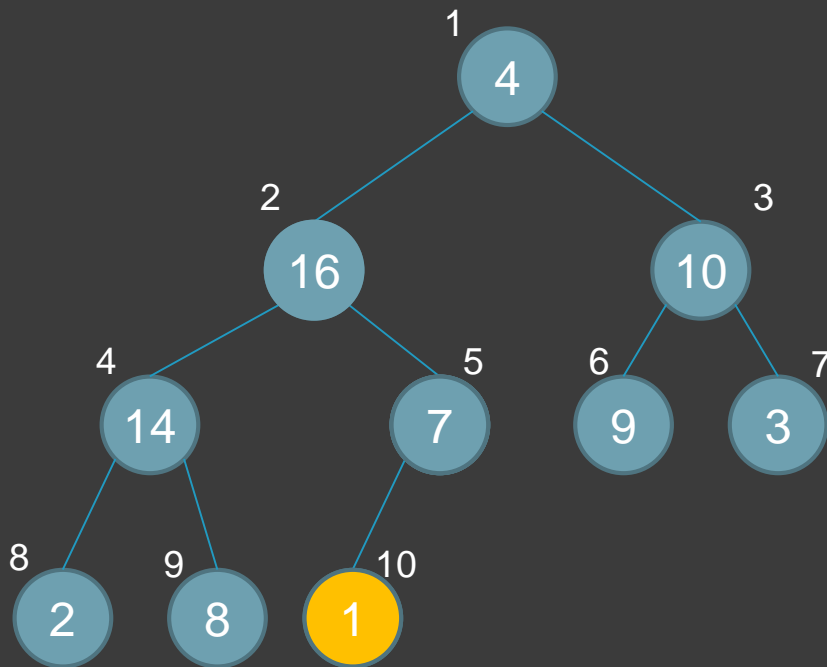
Then we MAX-HEAPIFY( $A, 7$ )  
 $A[7]$  is a leaf, so we're done here  
Go on to  $A[2]$



# Example

1	2	3	4	5	6	7	8	9	10
4	16	10	14	7	9	3	2	8	1

**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1  
  **do** MAX-HEAPIFY( $A, i$ )



$A[2]$  (1) is  $<$  its largest child (16), so  
swap  $A[2]$  and  $A[5]$

Then we MAX-HEAPIFY( $A, 5$ )  
 $A[5]$  is  $<$  its largest child (7), so  
swap  $A[5]$  and  $A[10]$

Then we MAX-HEAPIFY( $A, 10$ )  
 $A[10]$  is a leaf, so we're done

Go on to  $A[1]$

# Example

1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1

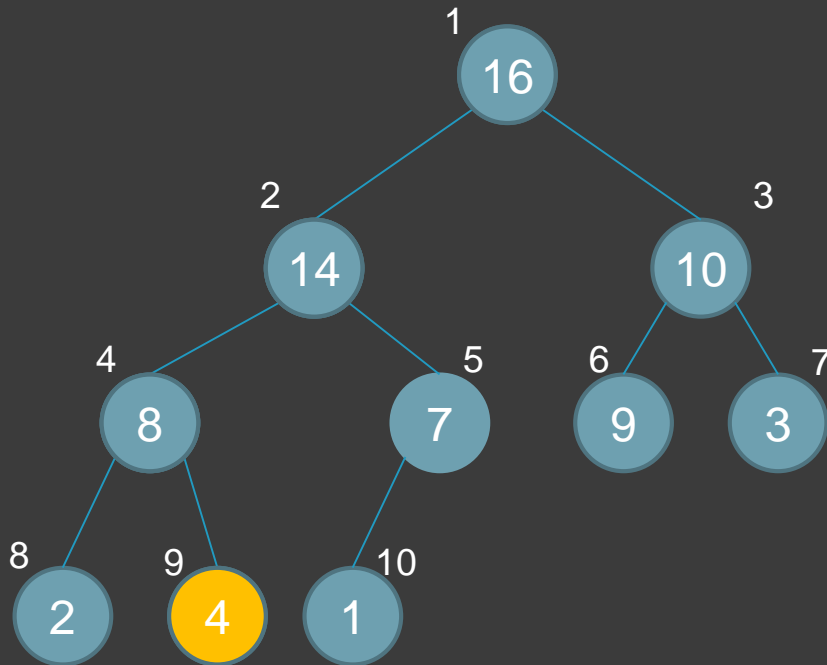
**for**  $i = \lfloor A.length/2 \rfloor$  **downto** 1  
**do** MAX-HEAPIFY( $A, i$ )

$A[1]$  (4) is  $<$  its largest child (16), so  
swap  $A[1]$  and  $A[2]$

Then we MAX-HEAPIFY( $A, 2$ )  
 $A[2]$  (4) is  $<$  its largest child (14),  
so swap  $A[2]$  and  $A[4]$

Then we MAX-HEAPIFY( $A, 4$ )  
 $A[4]$  (4)  $<$  its largest child (8), so  
swap  $A[4]$  and  $A[9]$  and then  
MAX-HEAPIFY( $A, 9$ )  $\rightarrow$  leaf

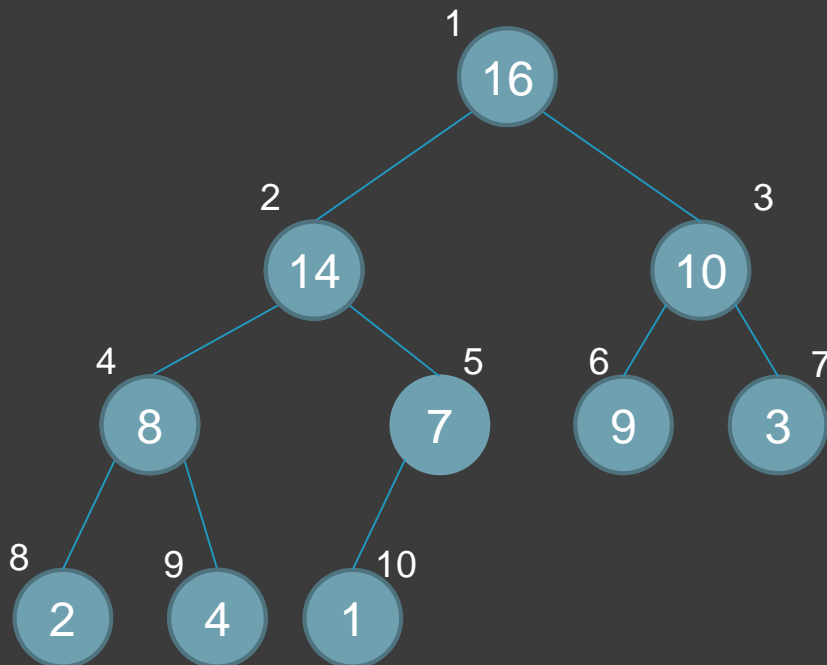
The heap is complete





# Example

1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1



Note:

We now have a valid Max-Heap

This does ***not*** mean the array is now sorted; it only means that the Max-Heap property is now valid at every location

But having a valid Max-Heap is the prerequisite for running Heapsort

# Heapsort (1)

- ⦿ Now that we can build a heap, and we can “re-heapify” a heap, we are ready to add one more simple step to create Heapsort.
- ⦿ We start with the array  $A$  in no particular order, and use BUILD-MAX-HEAP
- ⦿ At this point, with a proper heap,  $A[1]$ , the root of the heap, is the largest value in the array.

# Heapsort (2)

1	2	3	4	5	6	7	8	9	10
16	14	10	8	7	9	3	2	4	1

- Swap  $A[1]$  and  $A[n]$ . This puts the largest value in the  $n^{th}$  position, then  $A.heapsize--$

1	2	3	4	5	6	7	8	9	10
1	14	10	8	7	9	3	2	4	16

- That may have violated (probably did) the Max-Heap property at  $A[1]$ .
- Call MAX-HEAPIFY(1) to “re-heap” it
- That will put the next-largest item in  $A[1]$ .
- Swap  $A[1]$  and  $A[A.heapsize]$ , then  $A.heapsize--$
- Repeat...

# Heapsort Pseudocode

HEAPSORT(*A*)

1 BUILD-MAX-HEAP(*A*)

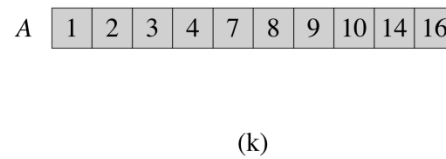
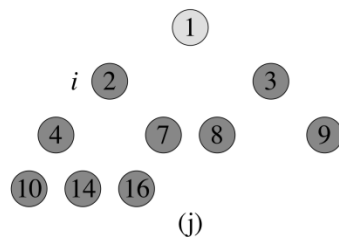
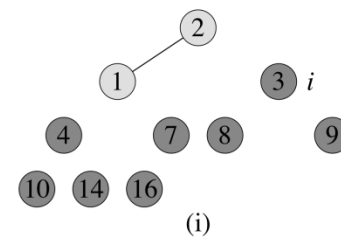
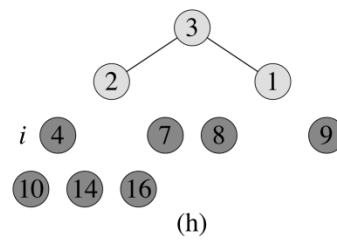
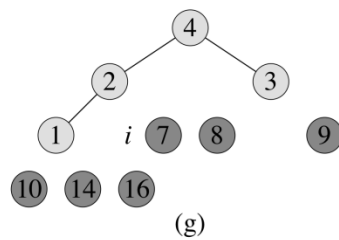
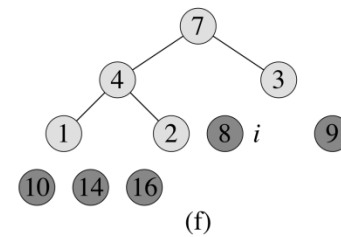
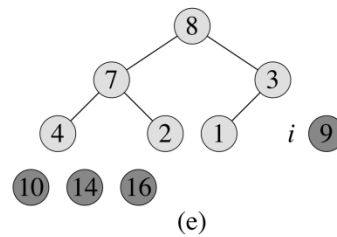
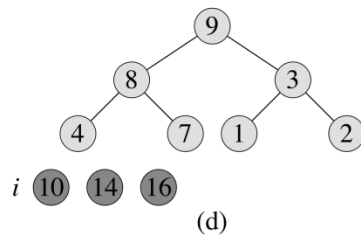
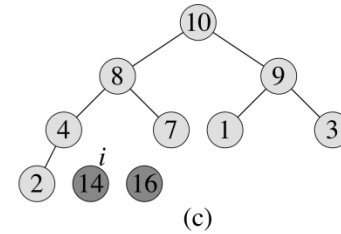
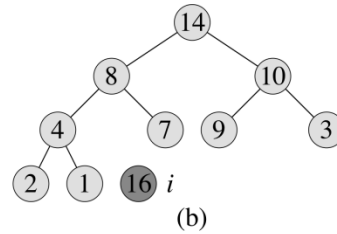
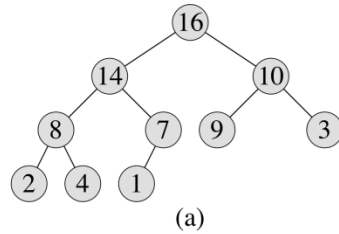
2 **for**  $i = A.length$  **downto** 2

3     exchange  $A[1] \leftrightarrow A[i]$

4      $A.heapsize = A.heapsize - 1$

5     MAX-HEAPIFY(*A*, 1)

# Example



# Heapsort - Summary

---

- ⦿ Combines best of Merge Sort...
  - $O(n \lg n)$  (guaranteed) execution time, and ...
- ⦿ ... and Insertion Sort
  - in-place sorting (requires only a constant amount of storage beyond the array to sort, unlike Merge Sort, which requires  $O(n)$  additional storage space)

# Proving BUILD-MAX-HEAP's Correctness

(See Chapter 2)

**Loop Invariant:** At the start of every iteration of the **for** loop, each node  $i + 1, i + 2, \dots, n$  is the root of a max-heap.

**Initialization:** We know that every node  $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$  is a leaf, which is the root of a trivial max-heap. Since  $i = \lfloor n/2 \rfloor$  before the first iteration of the **for** loop, the invariant is initially true.

# BUILD-MAX-HEAP's Correctness (2)

**Maintenance:** Children of node  $i$  are indexed higher than  $i$ , so by the loop invariant, they are both roots of max-heaps. Correctly assuming that  $i+1, i+2, \dots, n$  are all roots of max-heaps, MAX-HEAPIFY makes node  $i$  a max-heap root. Decrementing  $i$  reestablishes the loop invariant at each iteration.

**Termination:** When  $i = 0$ , the loop terminates. By the loop invariant, each node, notably node 1, is the root of a max-heap.



# Priority Queues

---

- ⦿ A data structure for maintaining a set of elements
- ⦿ Min-priority and max-priority queues (based on min-heaps and max-heaps, respectively).
- ⦿ The book focuses on max-priority queues; implementing min-priority queues is left as an exercise (6.5-3).

# Max-priority Queue Operations

- ⦿ A max-priority queue supports the following operations:
  - $\text{INSERT}(S, x)$  – inserts the element  $x$  into the set  $S$
  - $\text{MAXIMUM}(S)$  – returns the element of  $S$  with the largest key
  - $\text{EXTRACT-MAX}(S)$  removes and returns the element of  $S$  with the largest key
  - $\text{INCREASE-KEY}(S, x, k)$  – increases the value of element  $x$ 's key to the new value  $k$ , which is assumed to be at least as large as  $x$ 's current key value.

# Max-priority Queue vs. Min-priority

- Operations of max-priority and min-priority queues:

Max-priority	Min-priority
INSERT( $S, x$ )	INSERT( $S, x$ )
MAXIMUM( $S$ )	MINIMUM( $S$ )
EXTRACT-MAX( $S$ )	EXTRACT-MIN( $S$ )
INCREASE-KEY( $S, x, k$ )	DECREASE-KEY( $S, x, k$ )

# Max-priority Queue Application

---

- ⦿ Computer job scheduling
- ⦿ Max-priority queue keeps track of the jobs and their priorities
- ⦿ When one job finishes (or is interrupted), `EXTRACT-MAX` is used to select the highest-priority remaining job to run next
- ⦿ A new job can be added to the queue at any time with `INSERT`
- ⦿ The priority of a job already in the queue can be changed with `INCREASE-KEY`

# Min-priority Queue Application

---

- ⦿ Event-driven simulator
- ⦿ Items in the queue are events to simulate
- ⦿ Key values are event occurrence times
- ⦿ Events must be simulated in order of occurrence times
- ⦿ EXTRACT-MIN picks the next event to simulate
- ⦿ New events to be simulated are INSERTED into the queue.

# Implementing Priority Queues

---

- ⦿ Heaps can be used to implement priority queues
  - Often, we need to store a *handle* (pointer, index, etc.) to the corresponding application object in each heap element
  - How we implement and manage the handles will be application-dependent

# Implementing Max-priority Queue Operations

---

HEAP-MAXIMUM( $A$ )

1 **return**  $A[1]$

Runs in  $\Theta(1)$  time

# Implementing Max-priority Queue Operations

HEAP-EXTRACT-MAX( $A$ )

```
1  If  $A.heapsize < 1$ 
2      error “heap underflow”
3   $max = A[1]$ 
4   $A[1] = A[A.heapsize]$ 
5   $A.heapsize = A.heapsize - 1$ 
6  MAX-HEAPIFY( $A, 1$ )
7  return  $max$ 
```

Runs in  $\Theta(\lg n)$  time: MAX-HEAPIFY is  $\Theta(\lg n)$ ; all else is  $\Theta(1)$



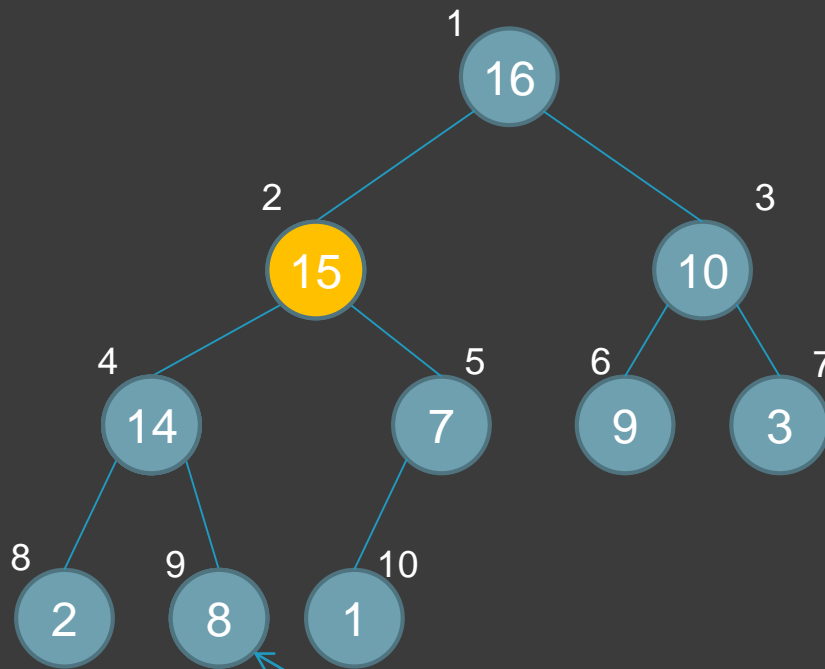
# Implementing Max-priority Queue Operations

HEAP-INCREASE-KEY( $A, i, key$ )

```
1  if  $key < A[i]$ 
2      error “new key is less than current key”
3   $A[i] = key$ 
4  while  $i > 1$  and  $A[PARENT(i)] < A[i]$ 
5      exchange  $A[i] \leftrightarrow A[PARENT(i)]$ 
6       $i = PARENT(i)$ 
```

Runs in  $O(\lg n)$  time: may move new value up heap to root; heap height is  $\Theta(\lg n)$

# HEAP-INCREASE-KEY Example



Swap this value with its parent as long as it is greater than its parent (or it has made it up to the root)

Increase  $A[9]$ 's key from 4 to 15

# Implementing Max-priority Queue Operations

MAX-HEAP-INSERT( $A, key$ )

1  $A.heapsize = A.heapsize + 1$  // Enlarge Heap

2  $A[A.heapsize] = -\infty$  // add new leaf

3 HEAP-INCREASE-KEY( $A, A.heapsize, key$ )

// Set desired priority (and percolate to proper location)

Runs in  $O(\lg n)$  time:

$$\Theta(1) + \Theta(1) + O(\lg n)$$

Line	1	2	3
------	---	---	---

# Summary

---

- Heaps can be used to implement priority queues with run time of  $O(\lg n)$  for any of the priority queue operations

## End of Chapter 6

---

Read your textbook  
for details.