

# ICSI 403

# DESIGN AND ANALYSIS OF ALGORITHMS

Lecture 11 – More Height-Balanced Trees:  
Red-Black Trees (Inserting)

# Review

- In BSTs (Binary Search Trees), the dynamic set operators (INSERT, DELETE, SEARCH, PREDECESSOR, SUCCESSOR, MAX, and MIN) all (can) take  $O(h)$  time, where  $h$  is the height of the tree.
- In a height balanced tree,  $h \approx \lg N$
- In a degenerate tree,  $h = N$
- We need to have a way to keep the tree balanced as we perform repeated insertions and deletions.

# A New Kind of Balanced Tree

- ◎ Red-Black Trees (1972 – Rudolph Bayer)
  - It's still a BST, and we keep it “mostly balanced”
  - Doesn't use “Balance Factor” like AVL
  - Instead, uses a single bit for “color” (“red” or “black”) and some very special rules.
    - By constraining the way nodes can be colored on any path from the root to a leaf, red-black trees (RBT's) ensure that no such path is more than twice as long as any other, so that the tree is always *approximately* balanced.

# Red-Black Trees (RBT's)

- Slightly different tree structure
- In the Huffman tree, data (characters) were only meaningful at a leaf.
- We're used to storing data in ANY node in a BST, and the only thing different about a leaf node was that it happened to not have any children (two NULL pointers).
- In an RBT, the leaves are ***always*** empty – our data resides in the internal nodes. Furthermore, all leaves are the SAME empty node, which Cormen calls *T.NIL*. The root's parent is also *T.NIL*

# Red-Black Trees (RBT's)

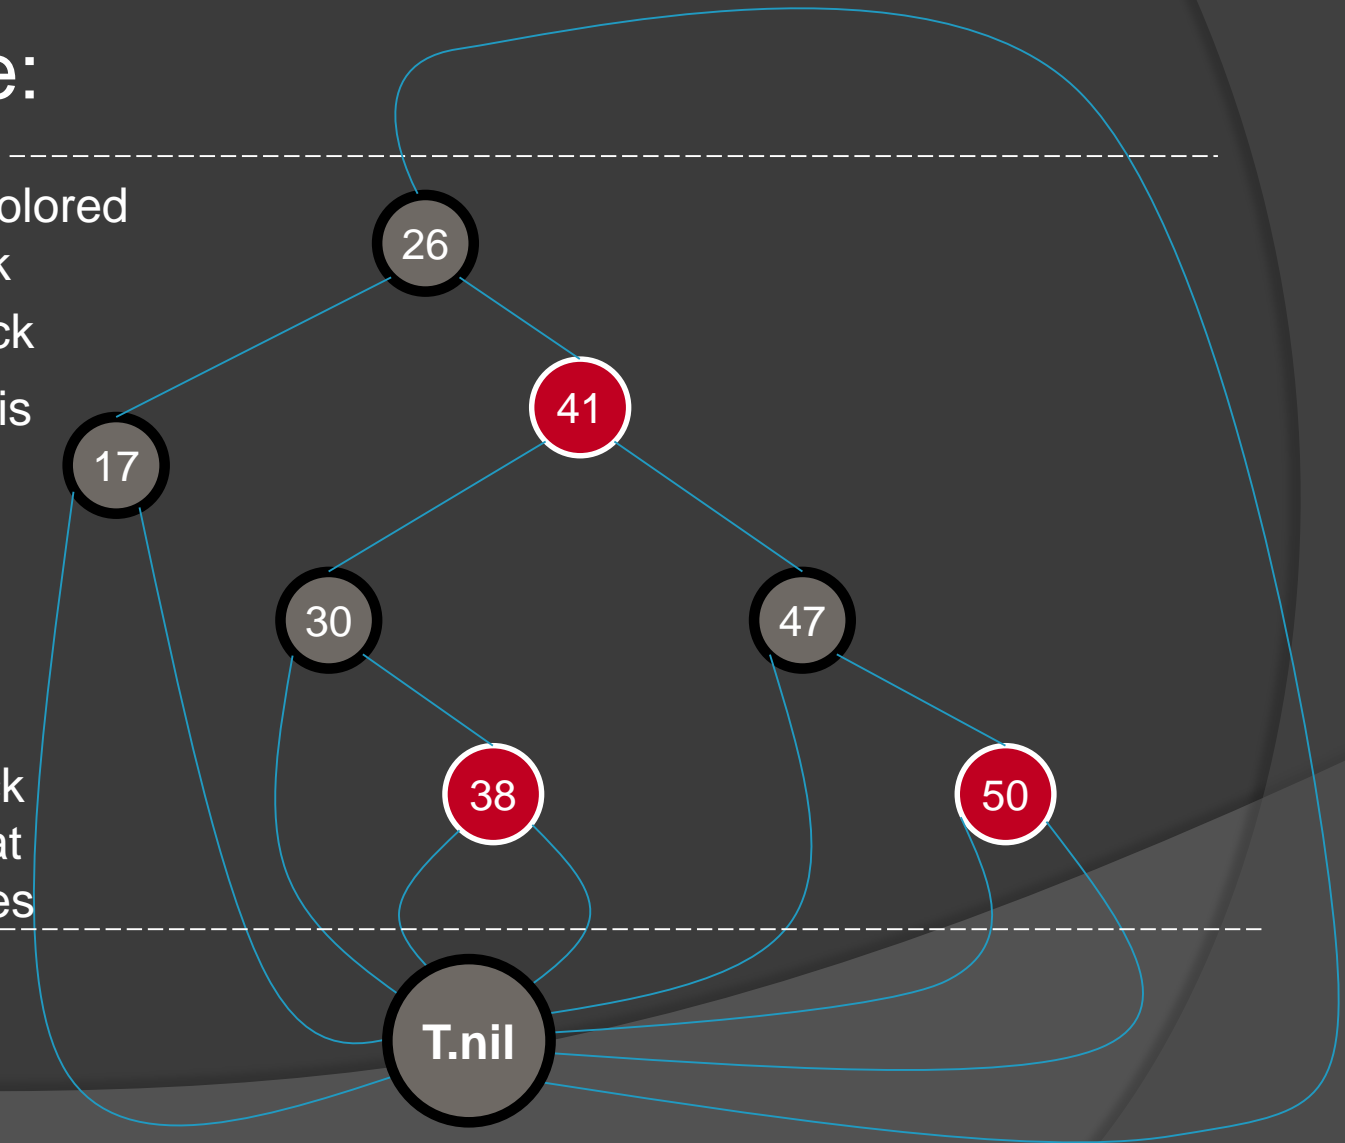
## ⦿ The “Ground Rules”:

- 1) Every node is colored either “Red” or “Black”
- 2) The root is black
- 3) Every leaf, *T.NIL*, is always black
- 4) If a node is red, then both of its children are black  
(hence, no two consecutive red nodes on a path from root to leaf)
- 5) For EVERY node, the number of black nodes between that node and the leaves is the same

# RBT Example

## ● Example:

- 1) Every Node is colored either Red or Black
- 2) The Root is Black
- 3) Every leaf, *T.nil* is *always* black
- 4) If a node is red, then **both** of its children are black
- 5) For every node, the number of black nodes between that node and the leaves is the same



# RBT Example

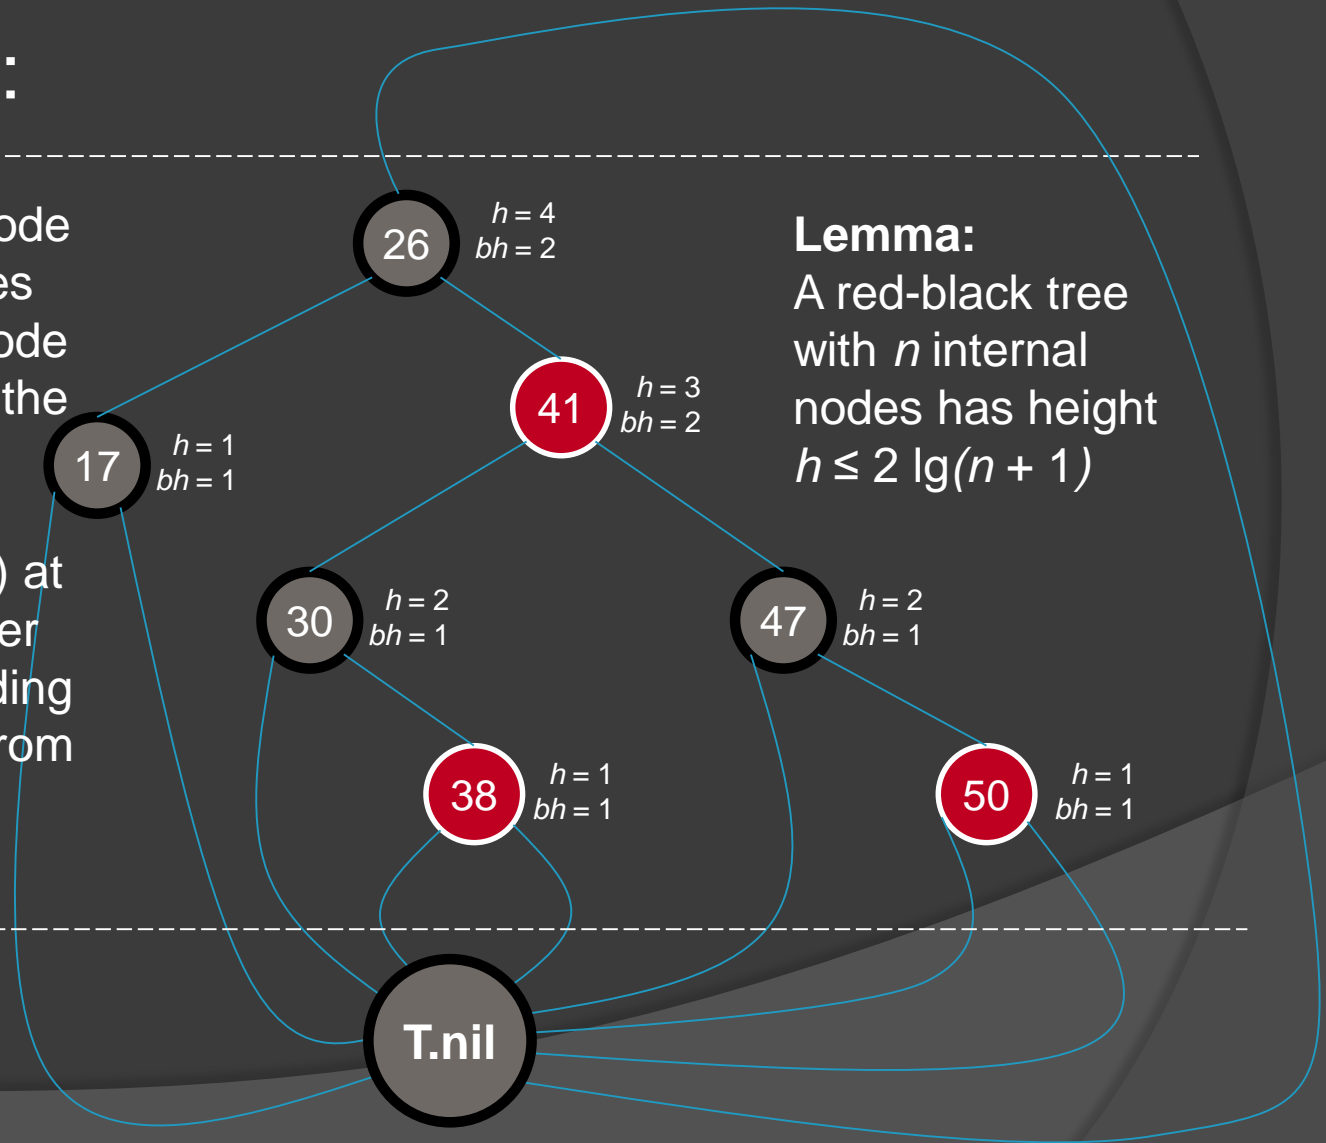
## Example:

The *height* ( $h$ ) at a node is the number of edges (links) between the node and the leaves along the longest path

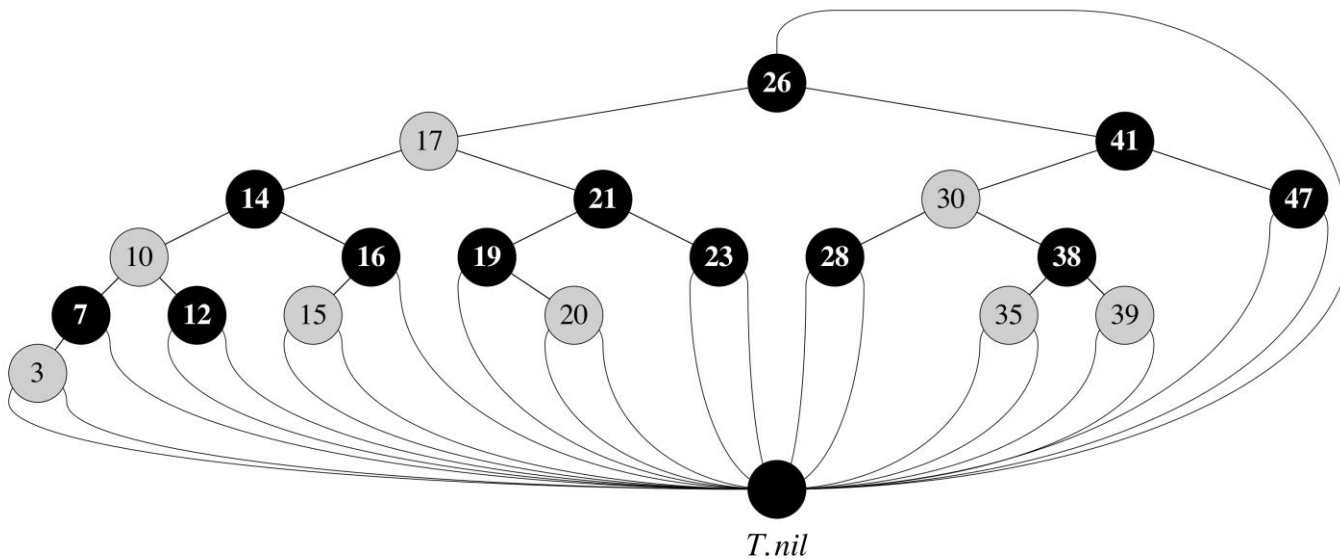
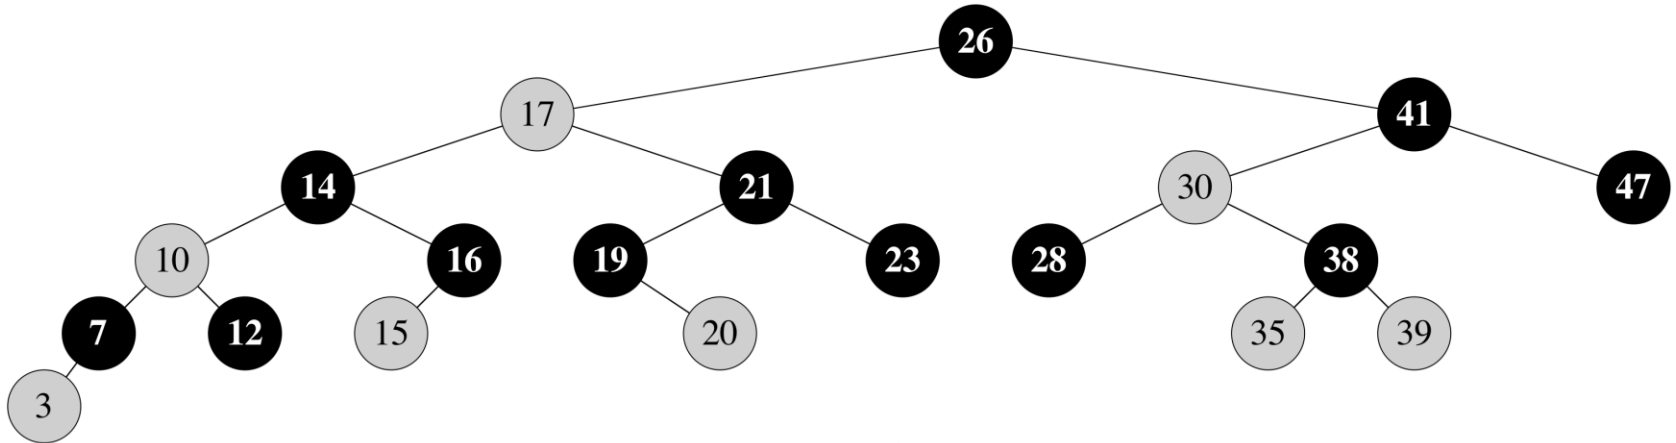
The *black height* ( $bh$ ) at a node  $X$  is the number of black nodes (including  $T.nil$ ) along the path from  $X$  to the leaves, not counting  $X$ .

### Lemma:

A red-black tree with  $n$  internal nodes has height  $h \leq 2 \lg(n + 1)$



# Example From the Text (1)





# Operations on RBTs

- ⦿ SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR work just like any other BST, and are all  $O(h) \rightarrow O(\lg n)$
- ⦿ INSERT and DELETE – not so easy.

# Operations on RBTs

- ◎ INSERT: What color to make the new node?
  - If we make it red, it might violate #4 (its parent might be red, and we can't have two successive red nodes)
    - 4) If a node is red, then **both** of its children are black
  - If we make it black, it might violate #5 (it might change the number of black nodes between some ancestor and the leaves)
    - 5) For every node, the number of black nodes between that node and the leaves is the same

# Operations on RBTs

- ◎ DELETE: What color was the node we removed?
  - If Red, it's OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Also, cannot cause a violation of property 2, since if the removed node was red, it could not have been the root.
  - If Black, could cause there to be two reds in a row (violating property 4), and can also cause a violation of property 5. Could also cause a violation of property 2, if the removed node was the root and its child – which becomes the new root – was red.

# RBT Operations

- ◎ The BST tree algorithms INSERT and DELETE run in  $O(\lg n)$  time, but do not guarantee that the modified binary search tree will be a red-black tree.
  - They give NO balance-related guarantees
- ◎ To restore these properties, we must change the colors of some of the nodes in the tree and also change the pointer structure.
- ◎ We change the pointer structure through ***rotation***

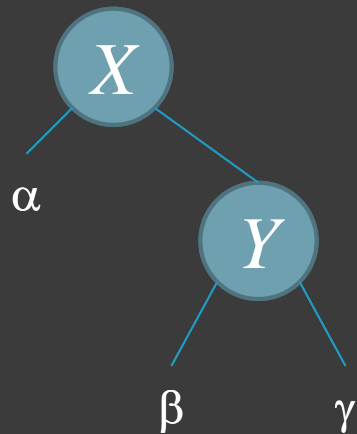
# RBT Rotations

- ⦿ Rotation is a local operation in a search tree that preserves the BST property (as in AVL).
- ⦿ When we do a left rotation on a node  $X$ , we assume that its right child  $Y$  is not  $T.nil$ ;  $X$  may be any node in the tree whose right child is not  $T.nil$ .
- ⦿ The left rotation “pivots” around the link from  $X$  to  $Y$ . It makes  $Y$  the new root of the subtree, with  $X$  as  $Y$ 's left child and  $Y$ 's left child as  $X$ 's right child.

# Rotations

- ⦿ The basic tree-restructuring operation
- ⦿ Needed to maintain red-black trees as balanced binary search trees
- ⦿ Changes the local pointer structure (only pointers are changed)
- ⦿ Won't upset the binary-search-tree property
- ⦿ There are both left and right rotations. They are inverses of each other (symmetric)
- ⦿ A rotation takes a red-black-tree and a node within the tree as arguments

# Left-Right Rotation



$$\alpha < X < \beta < Y < \gamma$$

$$\alpha < X < \beta < Y < \gamma$$

Note:  $\alpha$  stays the left child of  $X$ , and  $\gamma$  stays the right child of  $Y$   
Only  $X$ ,  $Y$ , and  $\beta$  change their relative positions

# Left-Rotate Pseudocode (p.313)

LEFT-ROTATE(T, x)

```
y = x.right           // y is x's right child
x.right = y.left       // Turn y's left subtree into
                        // x's right subtree
if y.left != T.nil
    y.left.p = x
y.p = x.p              // Link x's parent to y
if x.p == T.nil
    T.root = y
else if x == x.p.left
    x.p.left = y
    else x.p.right = y
y.left = x             // put x on y's left, which...
x.p = y                // ...makes x's parent be y
```

This all assumes  $T.root = T.nil$  and  $x.right \neq T.nil$



# Right-Rotate Pseudocode

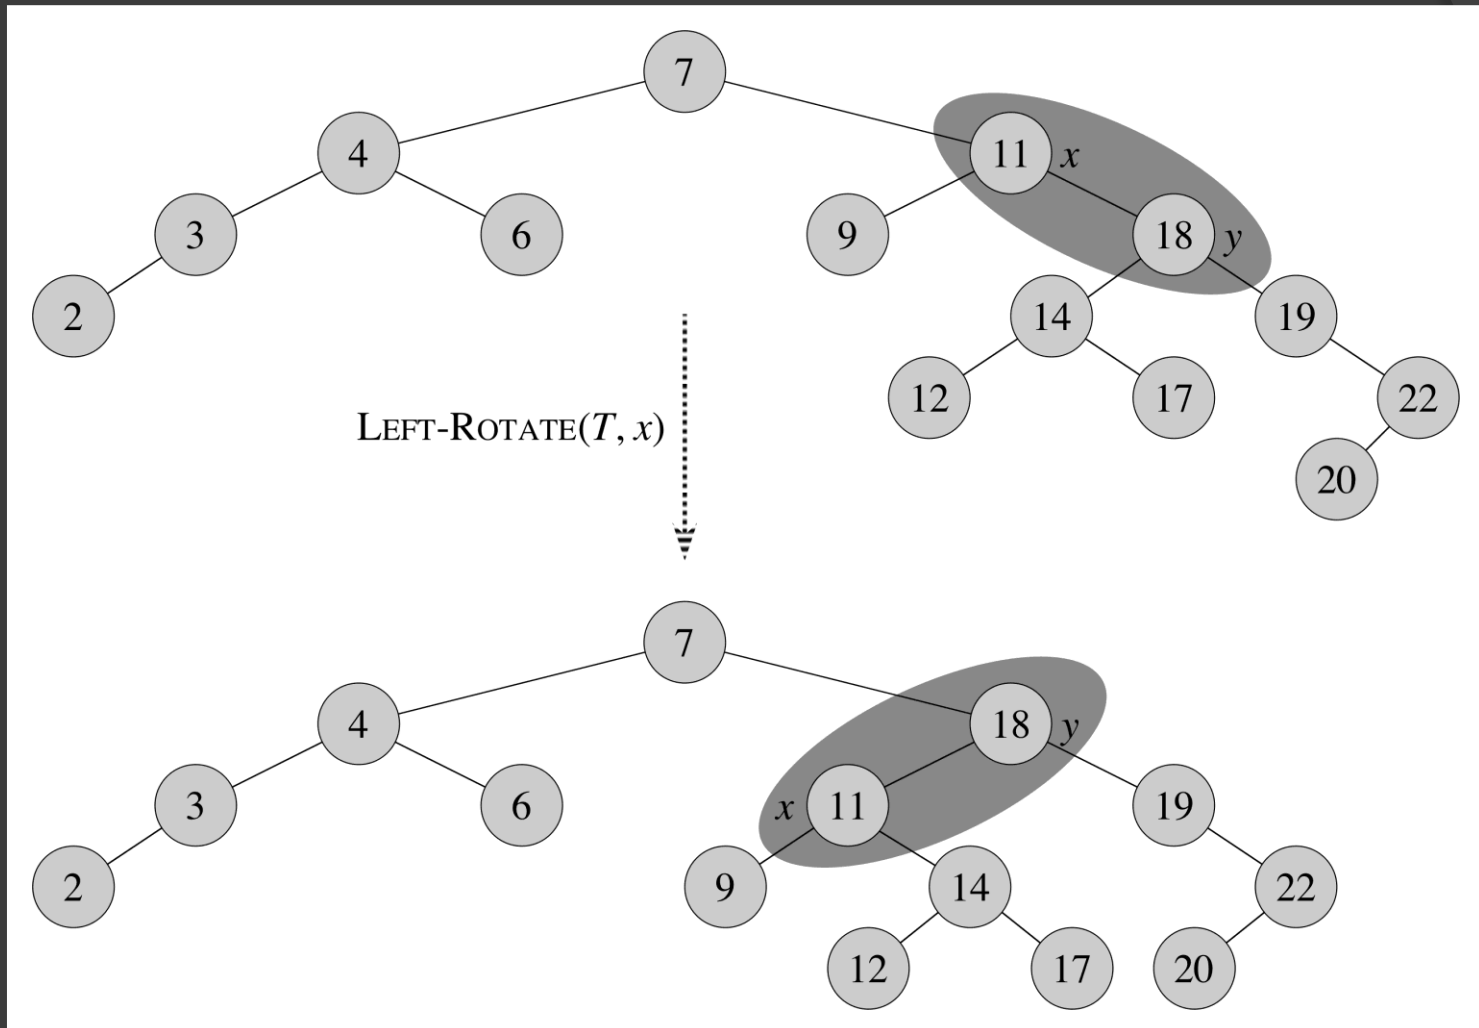
---

RIGHT-ROTATE(T, x)

// Everything is symmetric.

// Exchange all occurrences of “left” and “right”

# Left-Rotate Example



Rotation preserves the BST property

# Inserting Into a RBT

- ◎ Same (high-level) approach as AVL Insert:
  - Search the tree to find where this node belongs
    - Just like any “regular” BST insert
  - If the tree is empty, then this node becomes the root
    - Just like any “regular” BST insert
  - Otherwise, make the new node the appropriate child of the appropriate node
    - Just like any “regular” BST insert
  - Clean up any problems the insertion created
    - UNLIKE “regular” BST insert, but like AVL

# Inserting Into a RBT (p. 315)

```
RB-INSERT (T, z)      Insert node with key z into RBT T
    y = T.nil          x searches for insertion point
    x = T.root         y lags behind x (is its parent)
    while x != T.nil
        y = x
        if z.key < x.key x = x.left else x = x.right
    z.p = y
    if y == T.nil
        T.root = z
    else if z.key < y.key
        y.left = z
    else y.right = z
    z.left = T.nil     Our new node is at the bottom ...
    z.right = T.nil    .. of the tree, so its children are T.nil
    z.color = RED      Insert this node as a red one
    RBT-INSERT-FIXUP (T, z)  Fix anything we broke
```

# What Might We Have Messed Up?

- ◎ Which property might  $z$ 's insertion violate?
  1. *Every node is either red or black.* OK.
  2. *The root is black. If  $z$  is the root, then there's a violation. Otherwise, OK.*
  3. *Every leaf ( $T.nil$ ) is black.* OK.
  4. *If a node is red, then both its children are black.*  
*If  $z.p$  is red, there's a violation, because both  $z$  and  $z.p$  are red.*
  5. *For each node, all paths from the node to descendant leaves contain the same number of black nodes.* OK.

# What Might We Have Messed Up?

## ⦿ Notation for the next couple of slides:

- $z.p$  is  $z$ 's parent,  $z.p.p$  is  $z$ 's grandparent
- $y$  points to the “right uncle” of  $z$  (the right child of  $z$ 's grandparent)
- If there is a violation of the red-black properties, there is at most one violation, and it is of either property 2 or property 4.
  - If there is a violation of #2, it occurs because  $z$  is the root and is red.
  - If there is a violation of #4, it occurs because both  $z$  and  $p.z$  are red.

# What Might We Have Messed Up?

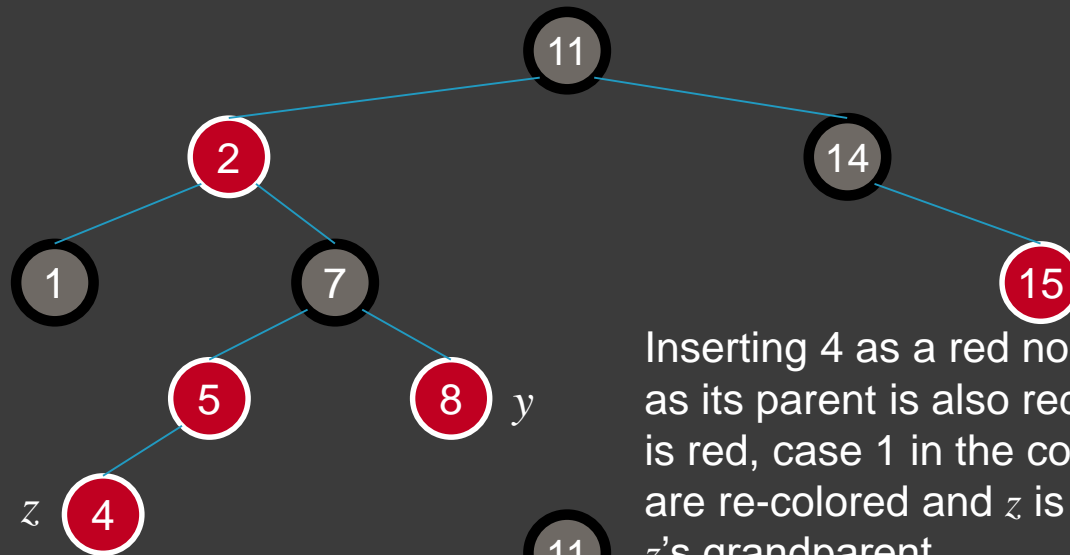
- ⦿ There are three cases (1, 2, 3)
- ⦿ Fixing one problem may create another,
- ⦿ If it does, however, it will not create any more than one problem (zero or one).
- ⦿ Eventually, we will solve a problem and not create a new one. At this point, we have re-balanced the tree.

# The Fixup Code (p. 316)

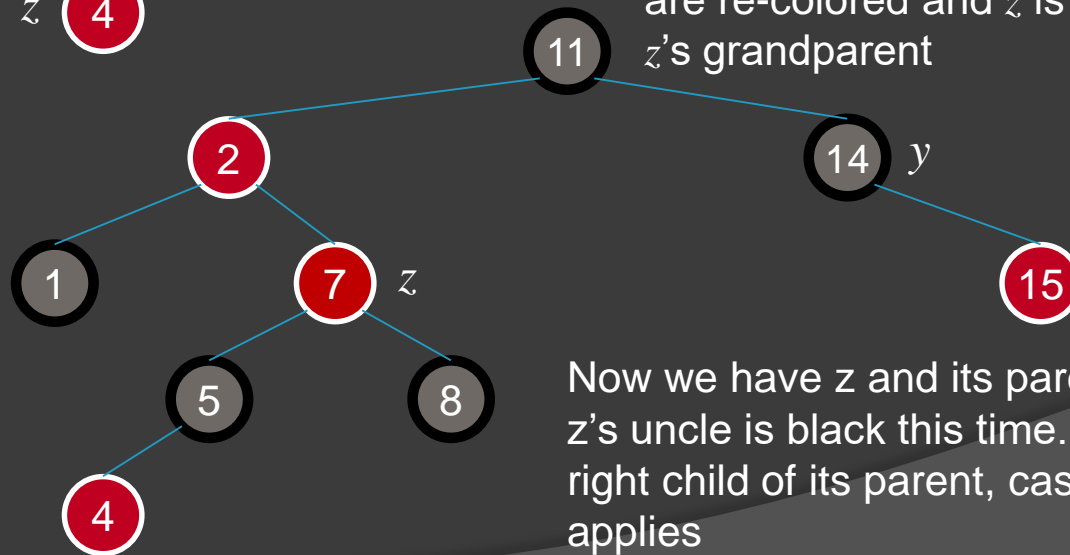
```
RBT-INSERT-FIXUP(T, z)
while z.p.color == red
    if z.p == z.p.p.left then
        y = z.p.p.right
        if y.color == RED then
            z.p.color = BLACK           Case 1
            y.color = BLACK             Case 1
            z.p.p.color = RED           Case 1
            z = z.p.p                   Case 1
        else if z == z.p.p.right
            z = z.p                     Case 2
            LEFT-ROTATE(T, z)           Case 2
            z.p.color = BLACK           Case 3
            z.p.p.color = RED           Case 3
            RIGHT-ROTATE(T, z.p.p)      Case 3
        else <“else” clause symmetric to “then”. Swap “left” and “right” >
T.root.color = BLACK
```



# Case 1

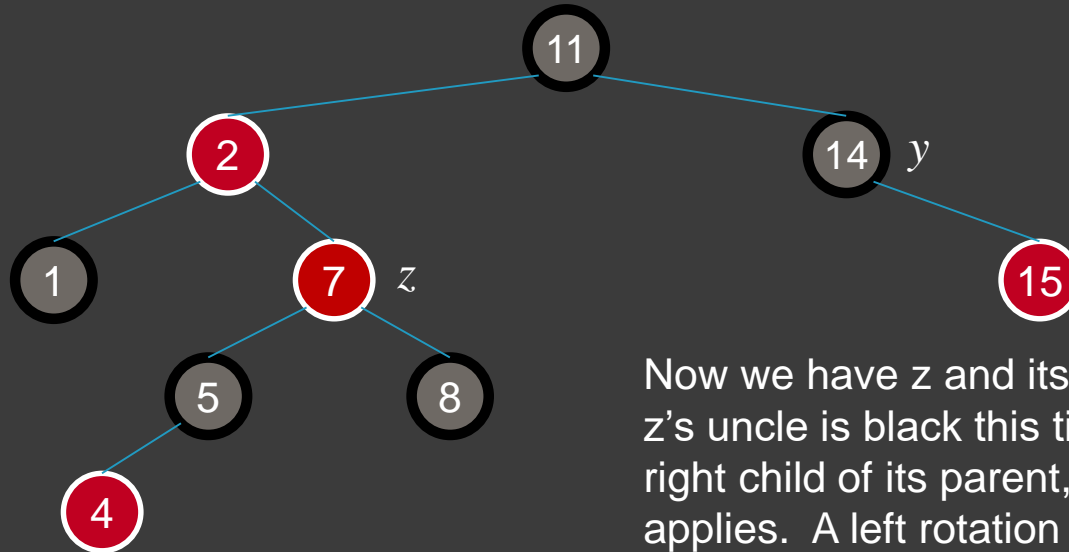


Inserting 4 as a red node ( $z$ ) violates rule 4, as its parent is also red. Since  $z$ 's uncle ( $y$ ) is red, case 1 in the code applies. Nodes are re-colored and  $z$  is moved up to point to  $z$ 's grandparent

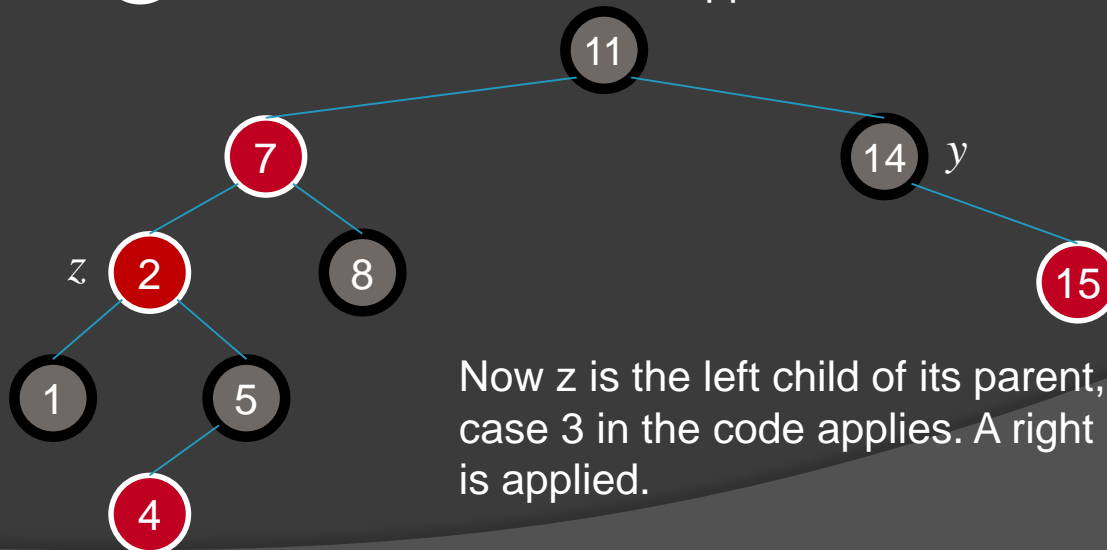


Now we have  $z$  and its parent both red, but  $z$ 's uncle is black this time. Since  $z$  is the right child of its parent, case 2 in the code applies

# Case 2

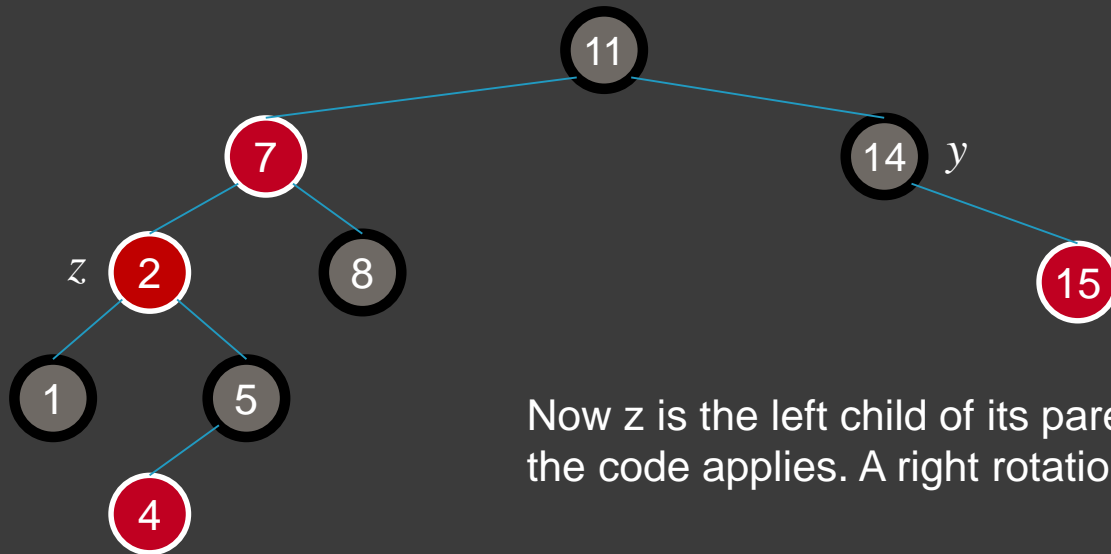


Now we have *z* and its parent both red, but *z*'s uncle is black this time. Since *z* is the right child of its parent, case 2 in the code applies. A left rotation is applied

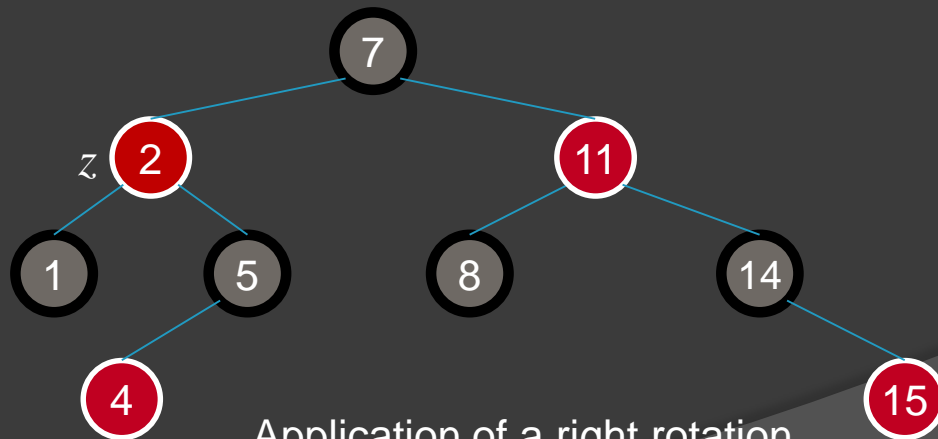


Now *z* is the left child of its parent, and case 3 in the code applies. A right rotation is applied.

# Case 3



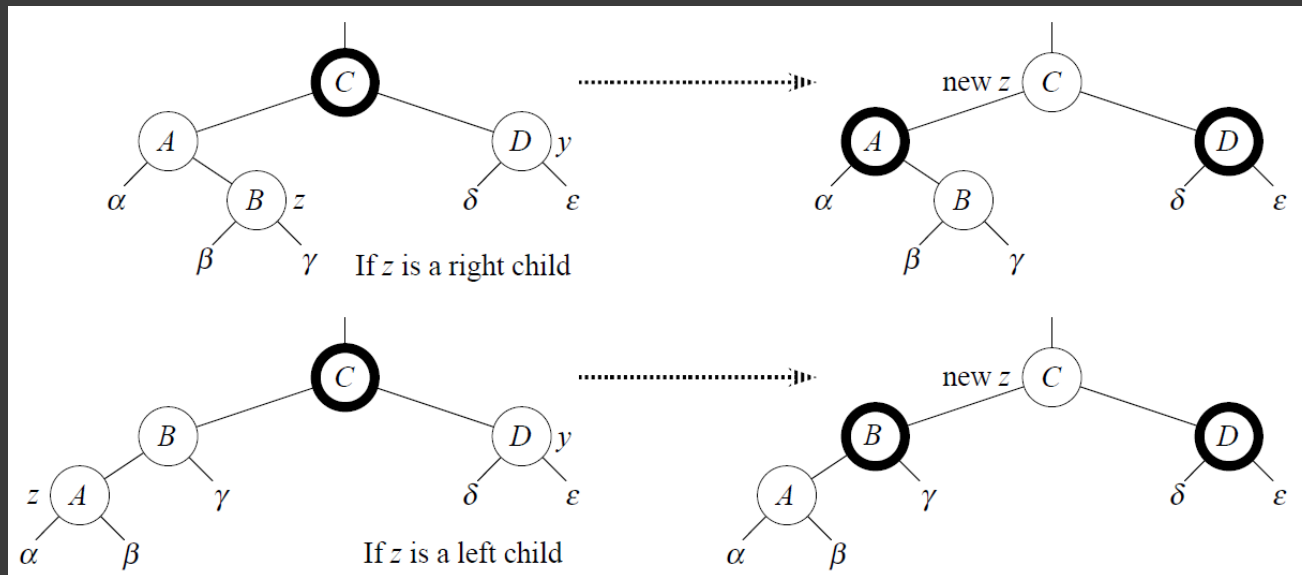
Now *z* is the left child of its parent, and case 3 in the code applies. A right rotation is applied.



Application of a right rotation results in a legal RBT

# The Three Cases – More Detail

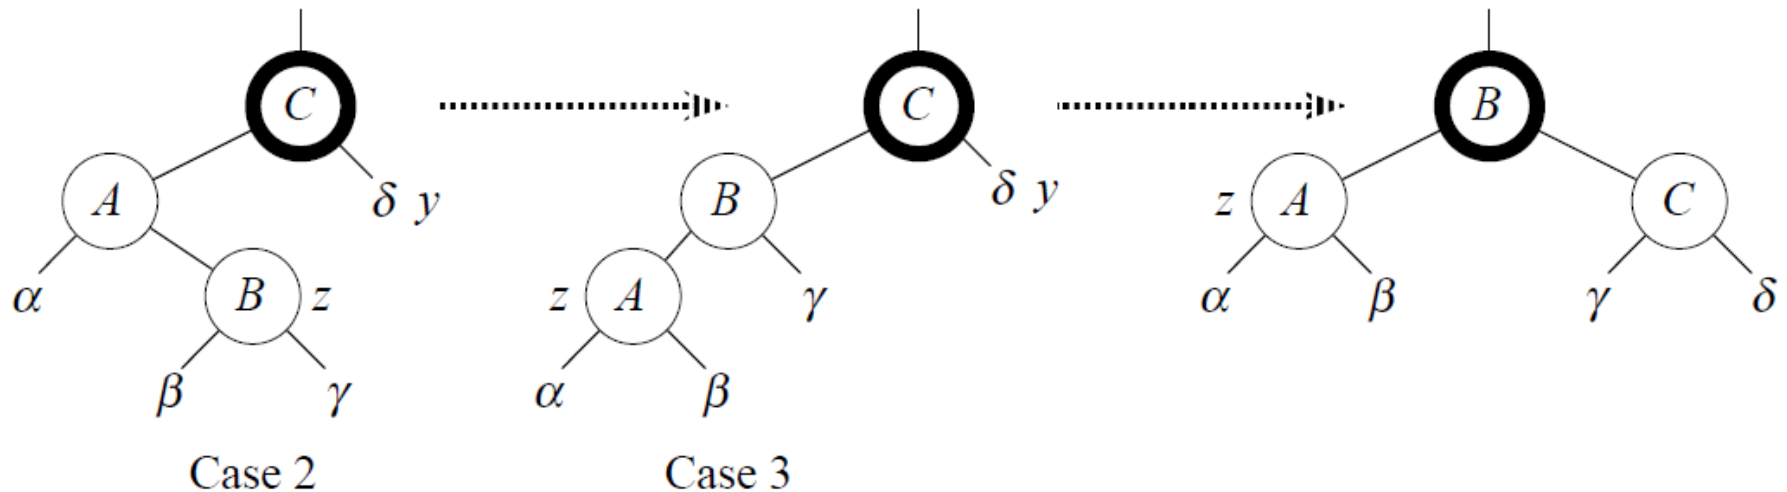
- Case 1:  $z$  is red,  $z.p$  is red, &  $z$ 's uncle  $y$  is red



- Since  $z.p.p$  is black, we can color both  $z.p$  and  $y$  black (fixing the problem of both  $z$  and  $z.p$  being red), and color  $z.p.p$  red (satisfying rule #5). Move  $z$  two levels up the tree, and go through the `while` loop again

# The Three Cases – More Detail

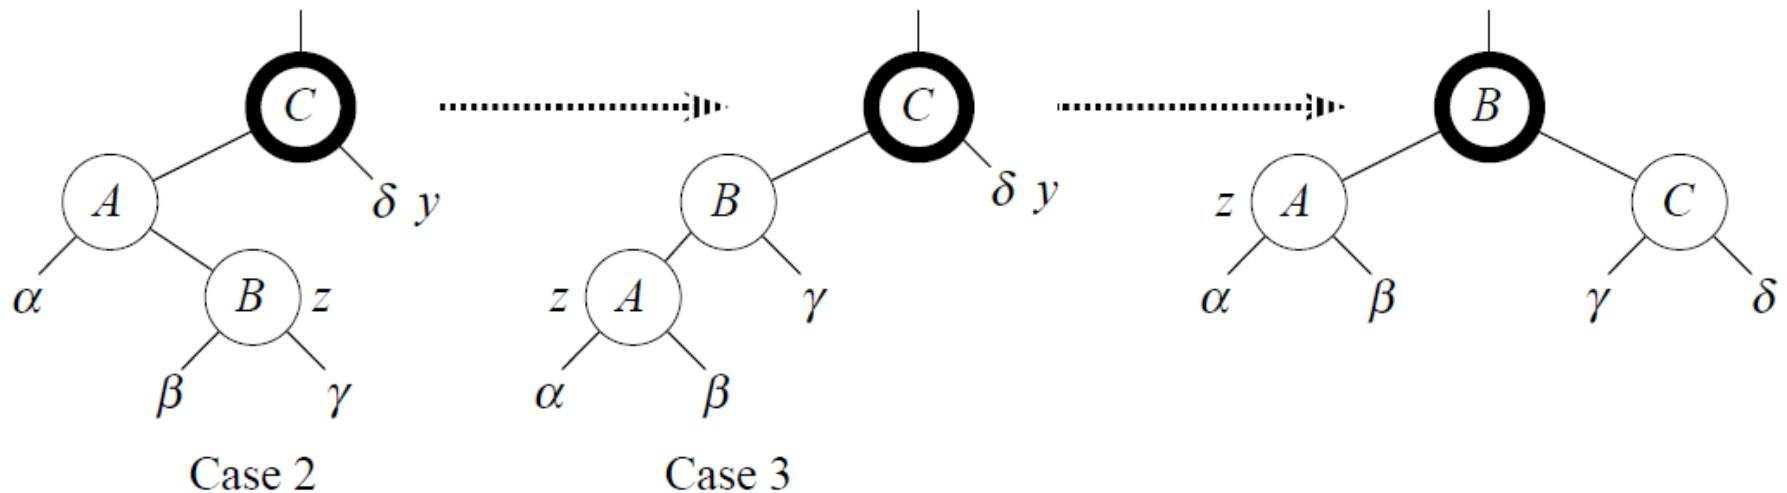
- Case 2:  $z$ 's uncle  $y$  is black,  $z$  is a right child



- Left rotate around  $z.p$ . Now  $z$  is a left child, and both  $z$  and  $z.p$  are red. Takes us immediately to case 3.

# The Three Cases – More Detail

- Case 3:  $z$ 's uncle  $y$  is black,  $z$  is a left child



- Make  $z.p$  black and  $z.p.p$  red
- Then right rotate on  $z.p.p$
- No longer have 2 reds in a row
- $z.p$  is now black  $\rightarrow$  no more iterations

# RBT Insertion Fixup - Summary

## ⦿ RBT Insertion – The Three cases:

- 1:  $z$  is red,  $z.p$  is red, &  $z$ 's uncle  $y$  is red
  - Re-color  $z$ ,  $z$ 's parent, and  $z$ 's uncle (no rotation)
- 2:  $z$ 's uncle  $y$  is black,  $z$  is a right child
  - Do a left rotation around  $z$ . Takes us immediately to
- 3:  $z$ 's uncle  $y$  is black,  $z$  is a left child
  - Recolor  $z$ 's parent and grandparent
  - Do a right rotation around  $z$ 's grandparent
- Move  $z$  up two levels ( $z \leftarrow z$ 's parent's parent), and try again as long as  $z$ 's parent is red

# RBT Analysis – Insert Operation

- ⊙  $O(\lg n)$  time to get through RBT-INSERT up to the call of RBT-INSERT-FIXUP
- ⊙ Within RBT-INSERT-FIXUP :
  - Each iteration takes  $O(1)$  time.
  - Each iteration is either the last one or it moves  $z$  up 2 levels.
  - $O(\lg n)$  levels.  $O(\lg n)$  time.
  - Also note that there are at most 2 rotations overall.
- ⊙ Thus, insertion into a red-black tree takes  $O(\lg n)$  time



# Red-Black Trees - Summary

## Operations on red-black-trees:

- SEARCH  $O(h)$
- PREDECESSOR  $O(h)$
- SUCCESSION  $O(h)$
- MINIMUM  $O(h)$
- MAXIMUM  $O(h)$
- INSERT  $O(h)$
- DELETE  $O(h)$

## Red-black-trees guarantee that the height of the tree will be $O(\lg n)$

# Next Time

---

- ⦿ RBT Deletion (pp. 323 ff.)
- ⦿ Similar to insertion:
  - Start with “regular” BST deletion
  - Then run a routine to repair anything the deletion did to violate the RBT rules.
    - THAT Fix-up has 4 cases.

? Questions ?

---

? Questions ?