

# CSI 403

# DESIGN AND ANALYSIS OF ALGORITHMS

Lecture 12 – Introduction to Height-Balanced Trees  
(AVL)

# Review: Trees

## ⦿ Trees (in general)

- Root
- Internal Nodes
- Leaves

## ⦿ Binary Trees

- Zero, one, or two children per node

## ⦿ Nodes consist of (at a minimum):

- *Some* data
- Left and Right child pointers
- The child pointers are NULL if the corresponding child nodes do not exist

# Full Trees


---

- When we first started talking about binary trees, we said that a full (or complete) binary tree had no NULL pointers on internal nodes (only leaves have NULL pointers)


# Full Trees and Tree Height



1 node:  $h = 1$



3 nodes:  $h = 2$



7 nodes:  $h = 3$

15 nodes:  $h = 4$

Completely full  
trees will have  
 $2^h - 1$  nodes.

# Full Trees and Tree Height

- In general, a full tree with  $h$  levels will have

$$2^{(h-1)} \leq n \leq 2^{(h)} - 1 \quad \text{nodes}$$


$h$	$n_{min}$	$n_{Max}$
1	1	1
2	2	3
3	4	7
4	8	15
5	16	31
6	32	63
7	64	127
8	128	255

Full trees are, by definition, always balanced.

# Adding Items To A Binary Tree

- ⦿ The resulting tree depends on the order in which the items were added!
- ⦿ If we insert 1, 2, 3, 4, and 5, *in that order*, then:
  - 1 becomes the root,
  - 2 is the right child of 1,
  - 3 is the right child of 2,
  - 4 is the right child of 3, and
  - 5 is the right child of 4:

# Result When Items Inserted In Order



Tree degenerates into a simple linked list – it's IMBALANCED!

Searching and inserting are now  $O(n)$ , rather than  $O(\lg n)$

We've lost our “each step discards (roughly) half of the tree” advantage


In a tree with a million nodes, a degenerate tree will require checking (on average) 500,000 nodes; a balanced tree would require checking about 10 nodes

# Balanced Trees

- ⦿ We can't always control the order in which data is inserted in the tree
- ⦿ The order the data is inserted determines the degree to which the tree will remain balanced
- ⦿ What we need is a way to accept insertions in any order, and still maintain a balanced tree.
- ⦿ Inserting a new node will create a new leaf, and may alter the tree's height → imbalance
- ⦿ If inserting a new node causes the tree to become imbalanced, we need to be able to re-balance it



# Rebalancing The Tree



Suppose we could “re-balance”  
the tree from this...

... to *this*

# Balanced Trees

- ⦿ There are over 100 types of balanced search trees
- ⦿ Among the more often used types are AVL trees, B-Trees, and Red-Black Trees
- ⦿ In a perfectly balanced tree containing  $N$  nodes, the height  $h \leq (\lg N)$
- ⦿ AVL Trees don't guarantee perfect balance, but they DO guarantee  $h \leq (\sim 1.4404 \lg N)$
- ⦿ Red-Black Trees (later) guarantee  $h \leq (2 \lg N)$

# AVL Trees

## ◎ AVL Trees (1962)

- Developed by Adelson-Velski and Landis
- Every insertion (and deletion) can result in imbalance.
- The Approach:
  - Go ahead and make the insertion
  - If an imbalance occurs, detect it and *re-balance*!
- Re-balance operations are called “rotations”

# AVL Tree Node Structure

- ⦿ Data
- ⦿ Left and right Child Pointers
- ⦿ Balance Factor (-1, 0, or +1)
  - Balance Factor (BF) at any node is the height of the node's LEFT sub-tree minus the height of the node's RIGHT sub-tree.
  - When insertion (or deletion) pushes a BF from -1 to -2 or from +1 to +2, we re-balance the tree
  - Tree is not maintained with perfect balance; however, it does stay “nearly balanced” or “balanced enough”


# Consider This Tree



- Each node shows Data:BF
- BF at a leaf is always zero
- Height of 7's left sub-tree = 1
- Height of 7's right sub-tree = 1
- $BF\ of\ 7 = (1 - 1) = 0$

# Let's Insert A Node


- Insert 3



- Update BFs changed by insertion


- The height of 7's left subtree (2) – the height of its right subtree (1) = +1
- The height of 5's left subtree (1) – the height of its right subtree (0) = +1
- The balance factors at 7 and 5 now become +1
- Nodes 3 and 9 are leaves, and by definition, always have a BF of 0

# Balance Factors Adjusted



# Let's Insert Again

- Insert 1



- Update BF's changed by insertion


- What will be the balance factors at each node after this insertion?



# BFs Updated

An imbalance occurs when an insertion pushes a balance factor from +1 to +2 or from -1 to -2

Imbalanced Nodes!



Rotations must do two things:

1. Reduce the sub-tree's height by 1
2. Preserve the BST Property (same in-order traversal order)

***Rotations are nothing more than child pointer changes!***

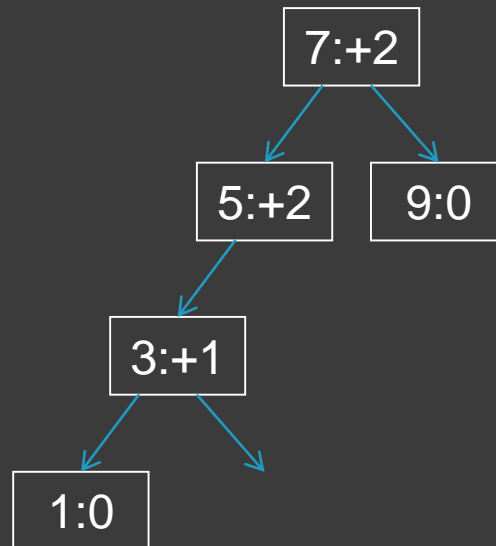
Relative to the node we just inserted, rotations occur below the most recent ancestor with a BF of +2 or -2

When an imbalance occurs, we solve the imbalance by applying a rotation.

# The Rotation Itself

Rotations must do two things:

1. Reduce the sub-tree's height by 1
2. Preserve the BST Property  
(same in-order traversal order)





Rearrange the nodes (change child pointers) to reduce the sub-tree's height

# The Rotation Itself


Rotations must do two things:

1. Reduce the sub-tree's height by 1
2. Preserve the BST Property  
(same in-order traversal order)



# The Rotation Itself

- Rotations must do two things:
1. Reduce the sub-tree's height by 1
  2. Preserve the BST Property  
(same in-order traversal order)




Finally, adjust BF's

What are the new (post-rotation) balance factors?

# The Completed Rotation


Rotations must do two things:

1. Reduce the sub-tree's height by 1
2. Preserve the BST Property  
(same in-order traversal order)




# Keep Inserting

Insert 11 and adjust BF's



# Keep Inserting


We could insert 10 without needing to re-balance, but inserting 13 would require a rotation.



Insert 13 and  
adjust BF's

After we insert 13, what will all of the BF's be?


# Keep Inserting



Rotate below most recent ancestor of inserted node that has a post-insertion BF of +2/-2




# Keep Inserting




Rotate below most recent ancestor of inserted node that has a post-insertion BF of +2/-2

# Keep Inserting



Re-attach rotated  
sub-tree


# Keep Inserting



Re-attach rotated  
sub-tree

Adjust Balance  
Factors


# Keep Inserting



# More On Insertions & Rotations

- ⦿ When we insert an item, and some node's BF goes to +2, we have two possibilities. The insertion was made in:
  - the left subtree of this node's left subtree (LL)
  - the right subtree of this node's left subtree (LR)
- ⦿ When we insert an item, and some node's BF goes to -2, we have two possibilities. The insertion was made in:
  - the left subtree of this node's right subtree (RL)
  - the right subtree of this node's right subtree (RR)


# Previous Example Revisited



- If we insert 1 or 2 (a left child of 3), then the BF at 3 goes to +1, and the BF at 5 goes to +2
- Relative to 5 (the node with the +2 BF), this is an insertion in the left subtree of the left subtree (LL)


- LL and RR are symmetrical
- If we insert 11 and then insert 13, that would be RR (the BFs at 9 and 11 would go to -2 and -1, respectively).

# Previous Example Revisited



- If we insert 4 (the right child of 3), then the BF at 3 goes to -1, and the BF at 5 goes to +2
- Relative to 5 (the node with the new +2 BF), this is an insertion in the right subtree of the left subtree (LR [go Left, then go Right])

# Previous Example Revisited



- If we insert 10 (the left child of 11), then the BF at 11 goes to +1, and the BF at 9 goes to -2
- Relative to 9 (the node with the new -2 BF), this is an insertion in the left subtree of the right subtree (RL [go Right, then go Left])



# Rotations

---

- ⦿ These four rotations (LL, RR, LR, and RL) are all we need.
- ⦿ The rotations we've done on previous slides have all been either LL or RR – the easy ones
- ⦿ LR and RL are more complicated – each has three sub-cases to check for

# A New Example

- Let's insert the names of the months (JAN through DEC) into an AVL tree (all comparisons are done alphabetically). We will see all four rotations in this exercise.
- For the sake of this exercise, we assume the names arrive in this order:  
MAR, MAY, NOV, AUG, APR, JAN,  
DEC, JUL, FEB, JUN, OCT, SEP.
- In each case, we will show the tree AFTER the insertion and BF adjustments

# Insert MAR (Into Empty Tree)


MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



No rotation  
Needed

# Insert MAY


MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



No rotation  
Needed


# Insert NOV

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



# Insert AUG

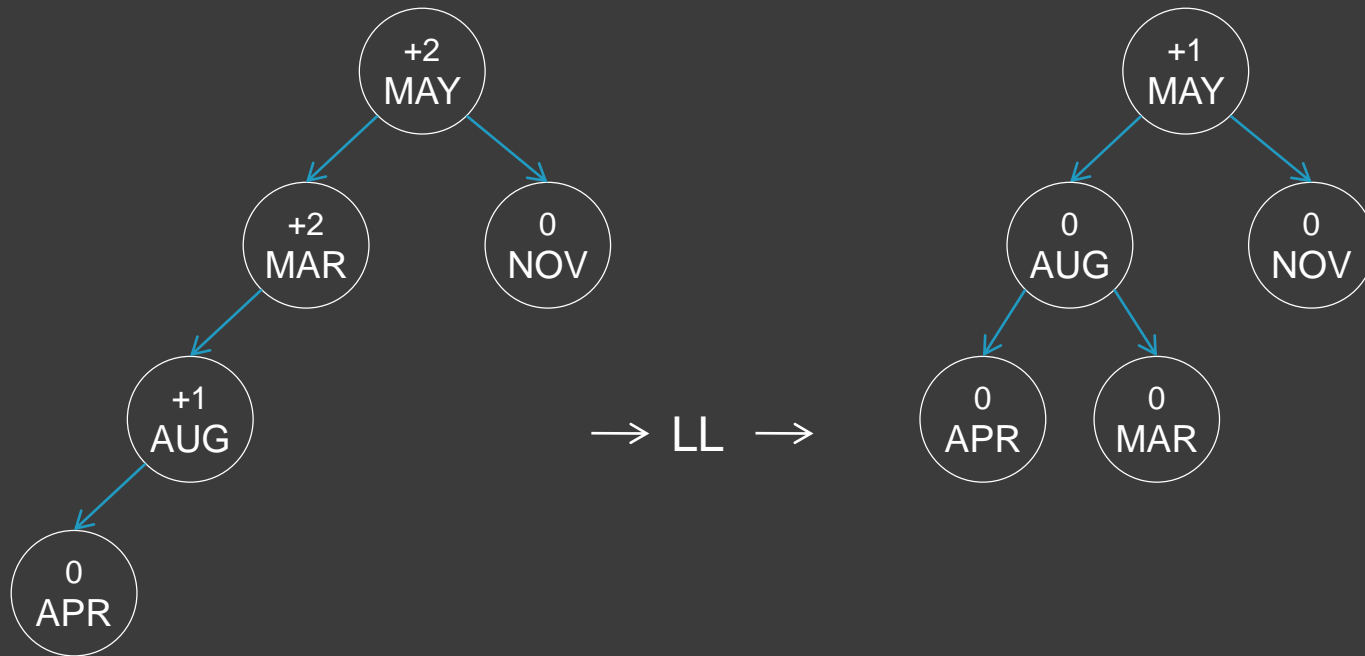
MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



No rotation  
Needed

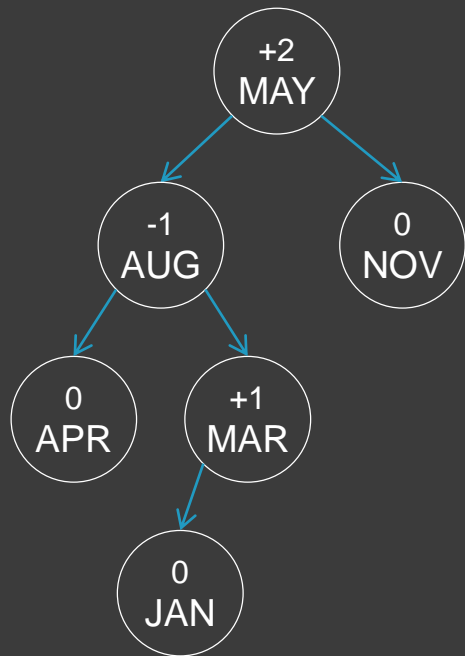
# Insert APR

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP

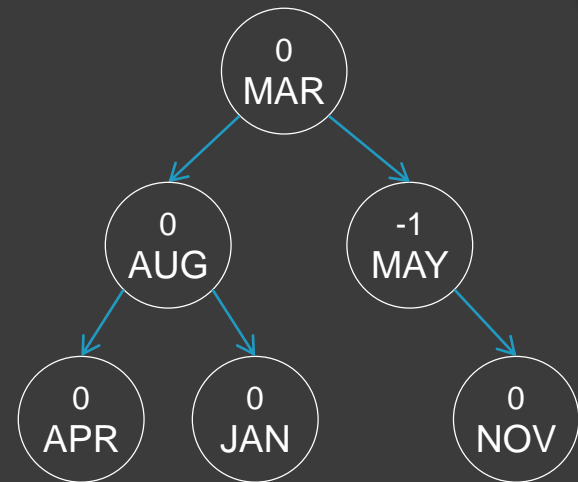


# Insert JAN

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



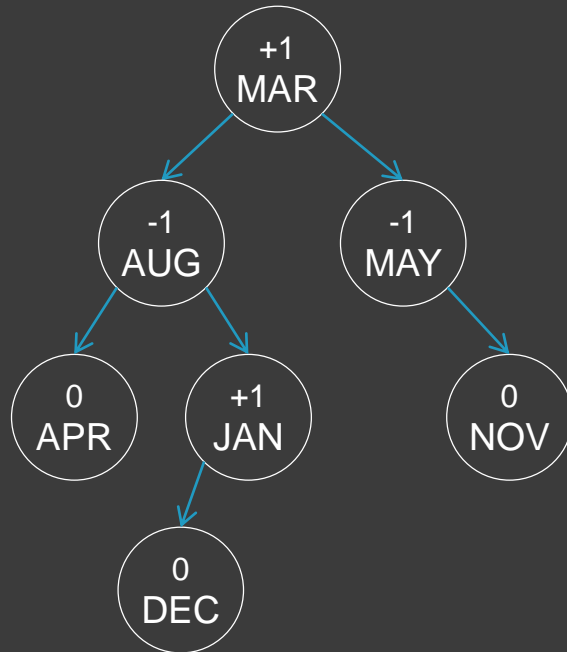
→ LR →





# Insert DEC

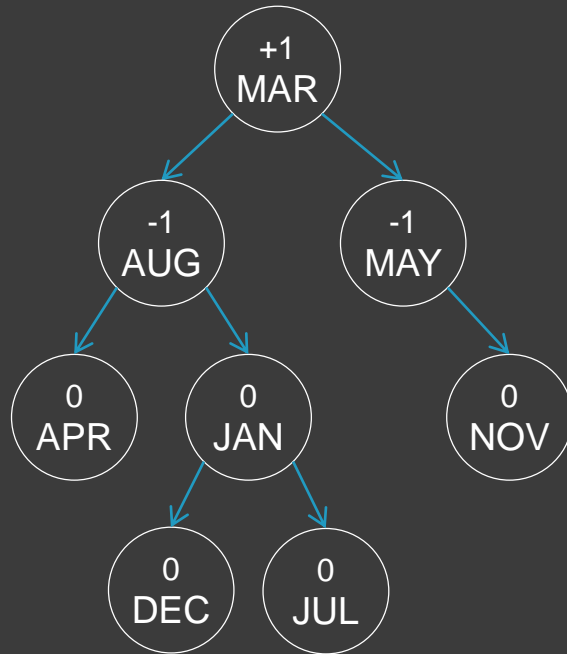
MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



No rotation  
Needed

# Insert JUL

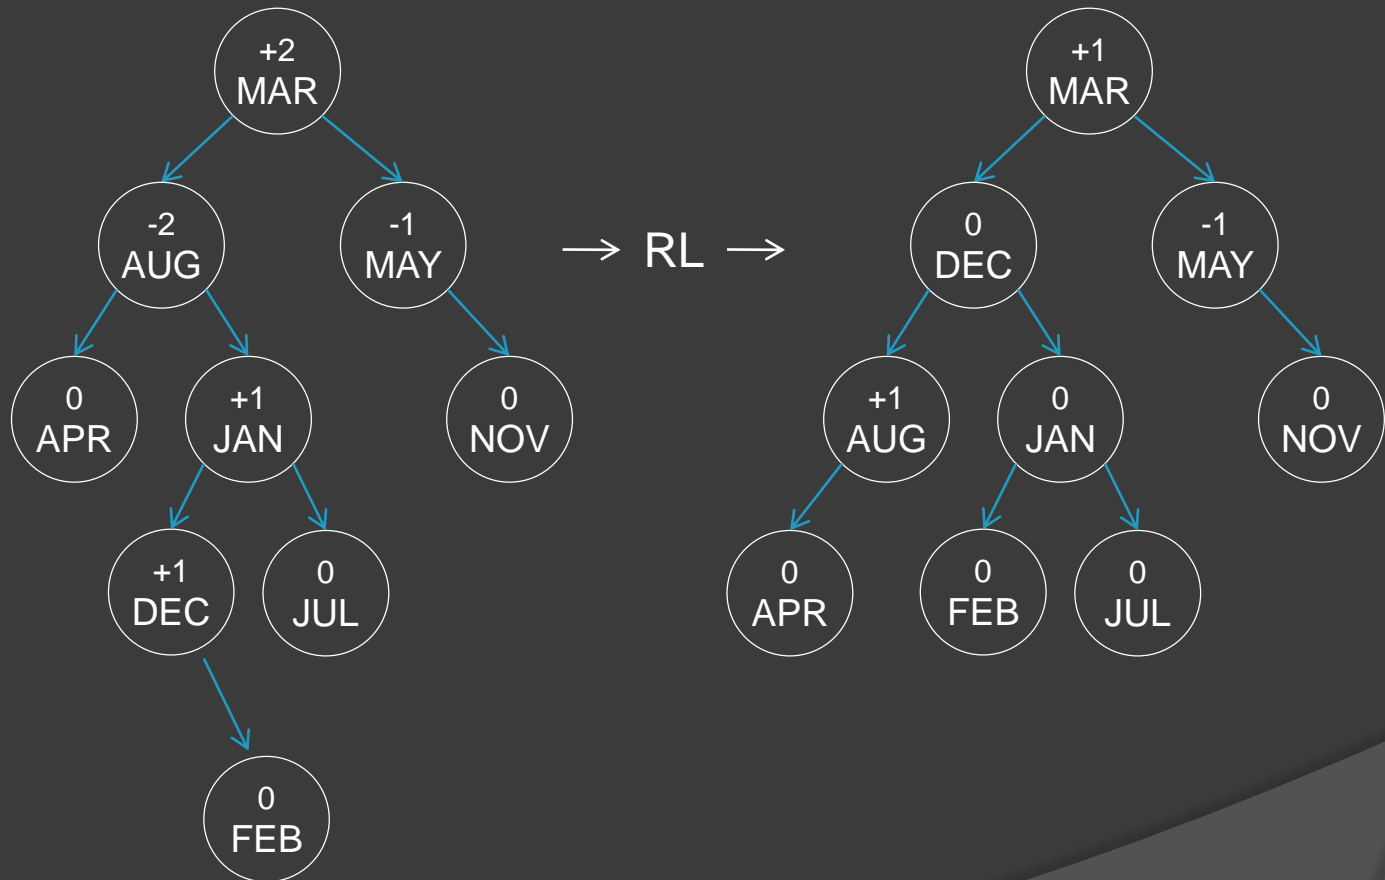
MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



No rotation  
Needed

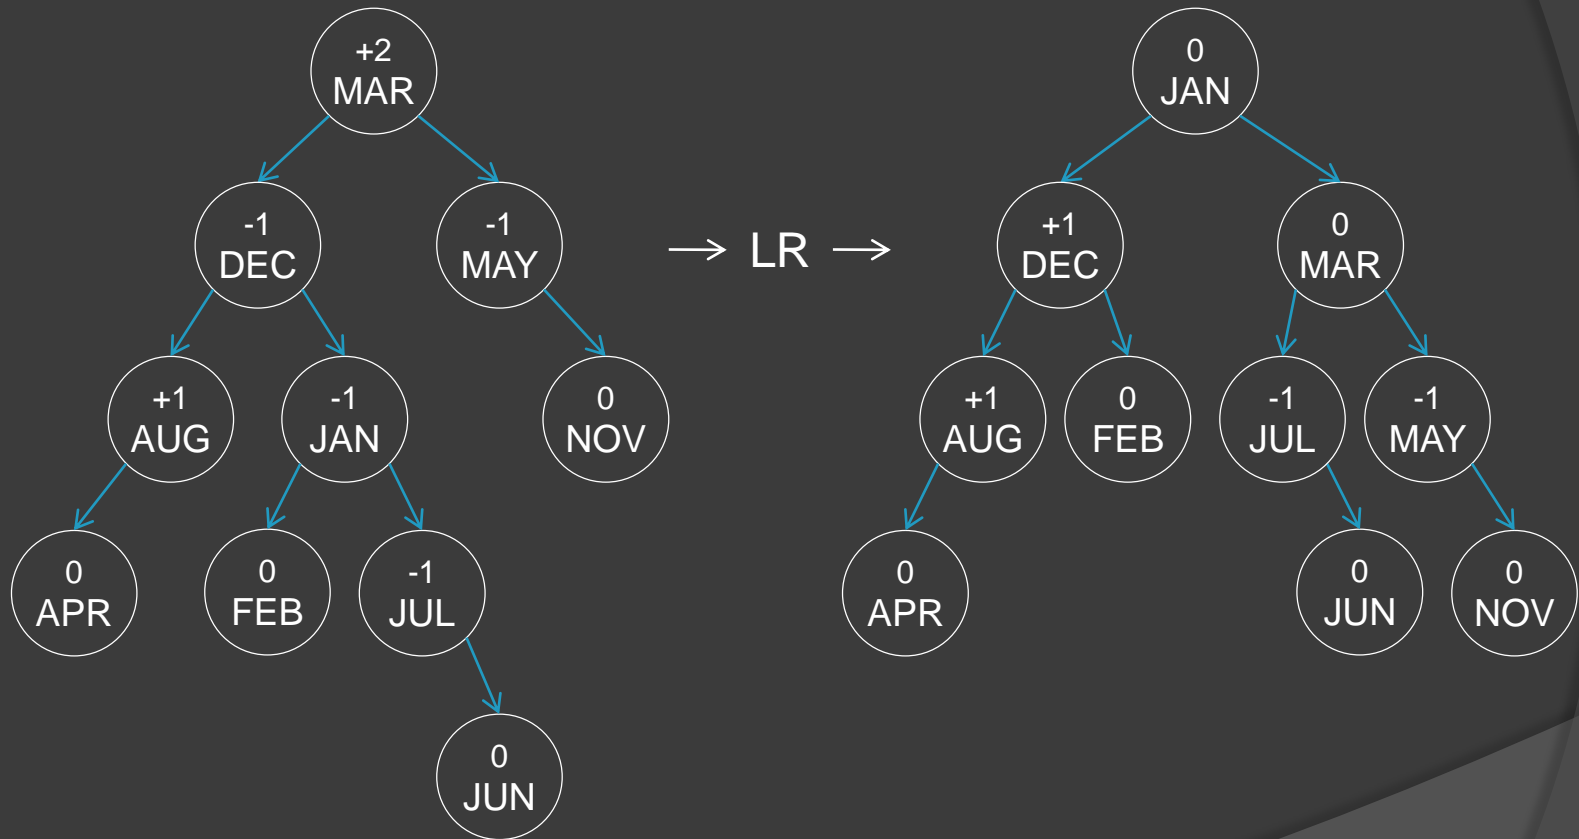
# Insert FEB

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



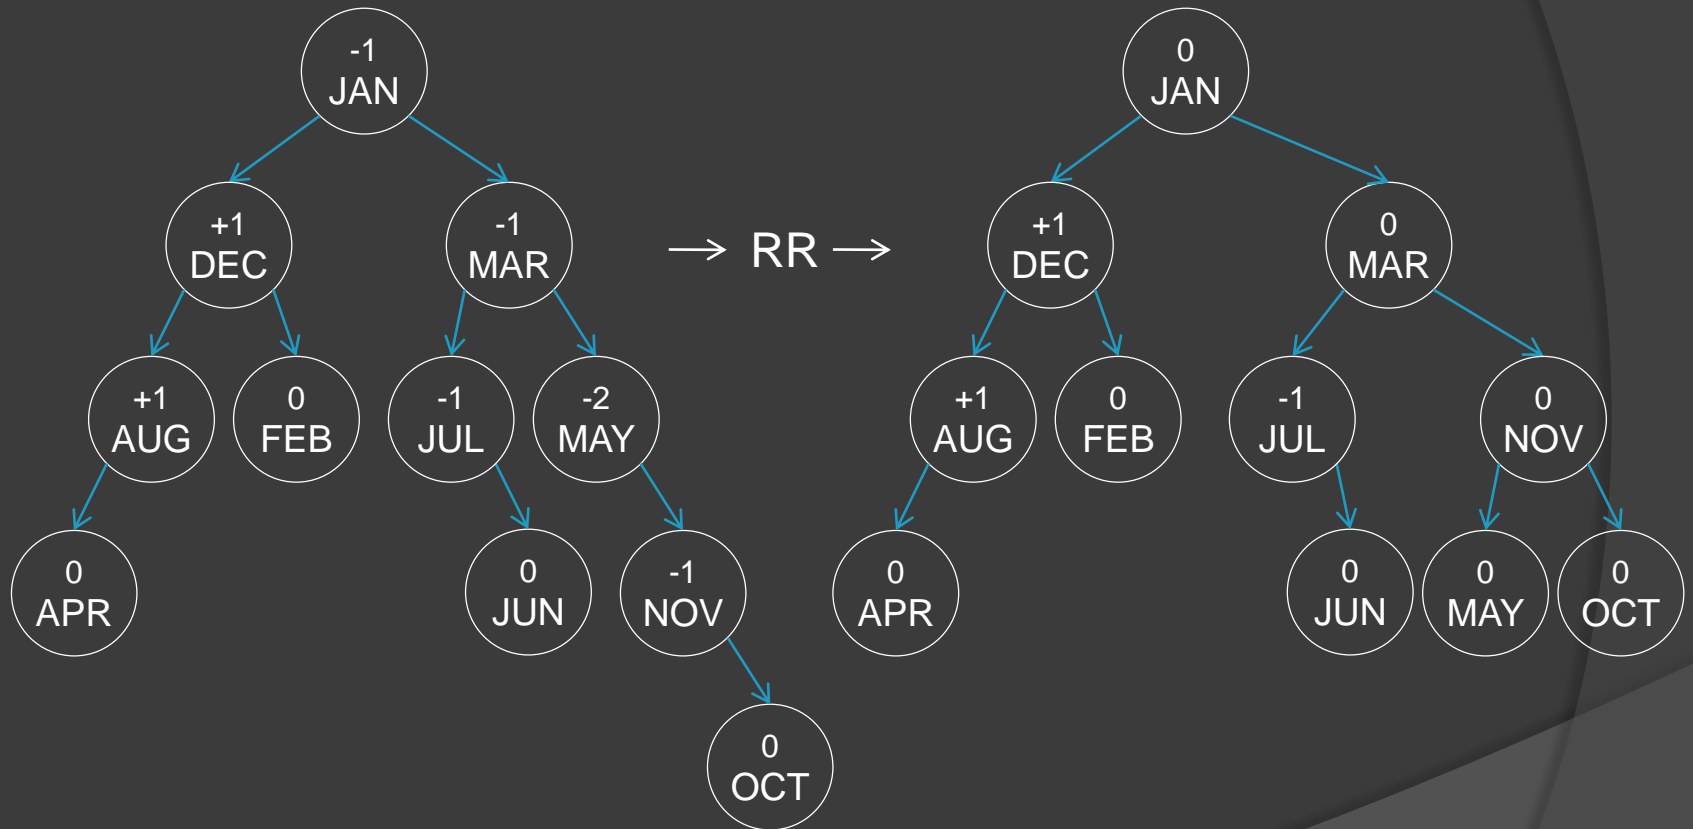
# Insert JUN

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



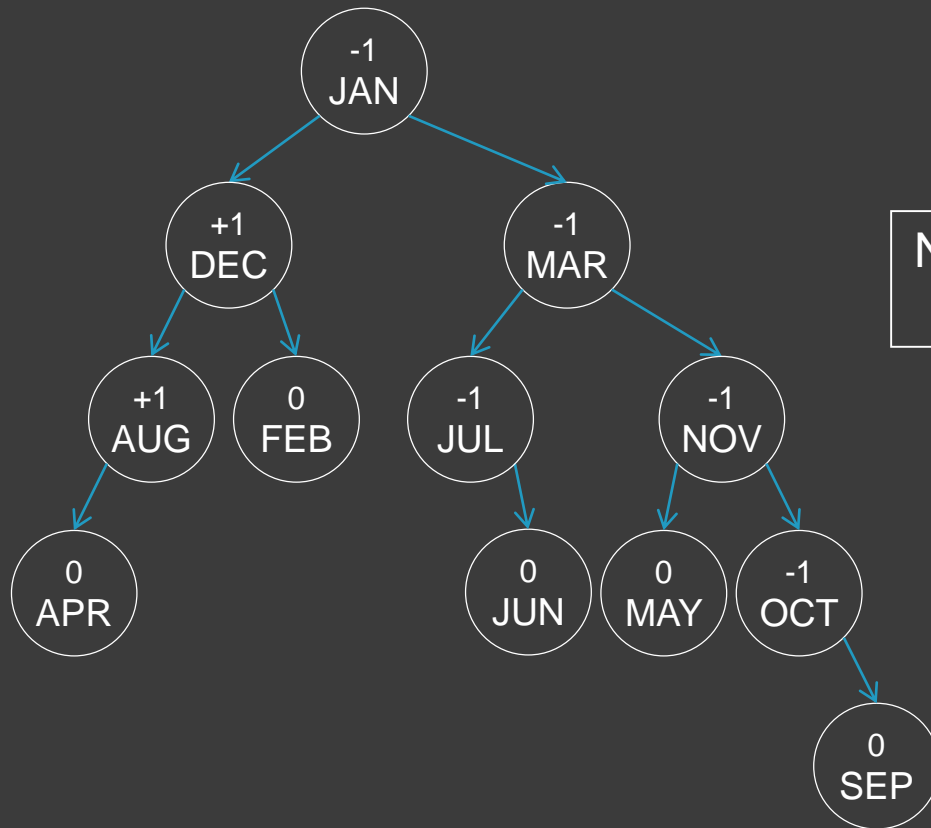
# Insert OCT

MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP



# Insert SEP


MAR  
MAY  
NOV  
AUG  
APR  
JAN  
DEC  
JUL  
FEB  
JUN  
OCT  
SEP




No rotation  
Needed

# The Rotations, Schematically (1)

Balanced subtree



Unbalanced following insertion






Height of  $B_L$  increases to  $h+1$

rotation type


LL →

Rebalanced subtree



Height of  $B_R$  increases to  $h+1$


RR →




Height of subtrees of B remain  $h+1$

# The Rotations, Schematically (2)

Balanced subtree





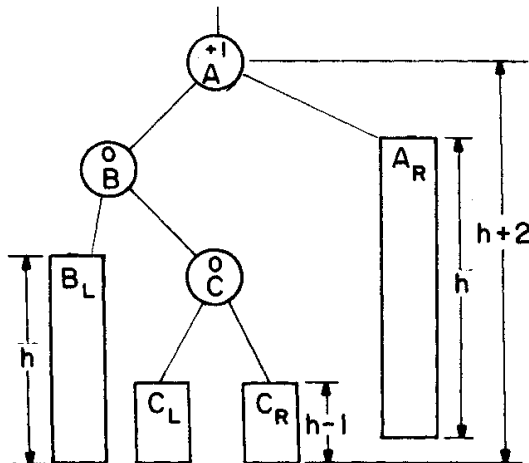

Unbalanced following insertion



rotation type

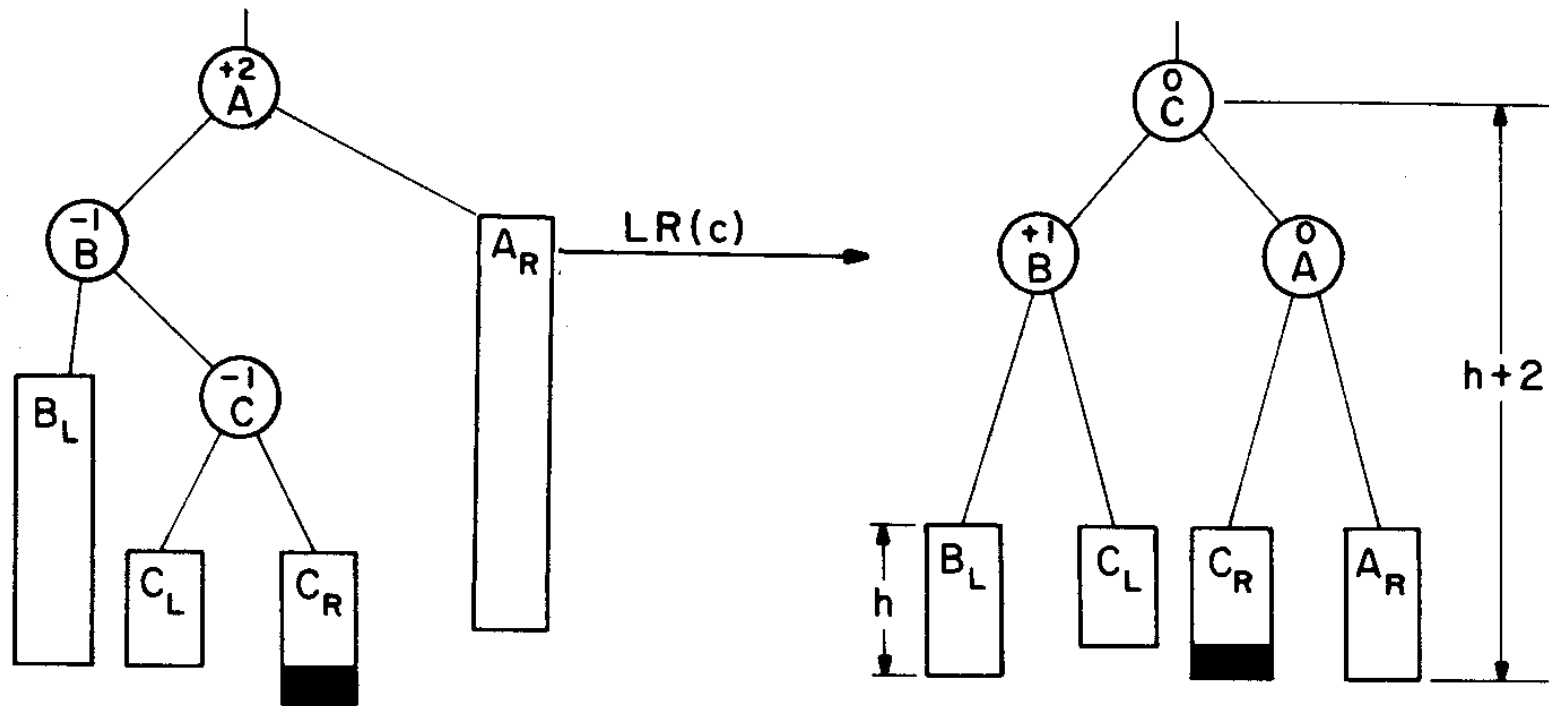
LR(a) →

Rebalanced subtree





# The Rotations, Schematically (3)



# The AVL Insert Process (Pseudocode)

## SUPER high-level:

1. If tree is empty, make new root node ( $BF=0$ )
2. Otherwise, scan through tree, looking for where what we want to insert belongs. If we find it (already in the tree), exit.
3. Make a new node, and make this new node the appropriate child of its new parent.
4. Adjust Balance Factors above new node
5. If imbalance occurs, rotate to rebalance.

# The AVL Insert Process (In C): 1/8

```
node *root=NULL;
void AVL_Insert(data X) {
    node *Y;           // The new node we insert
    node *A, *B, *F;   // see below
    node *C, *CL, *CR  // for description
    int d;              // Used to adjust BF's
    if (root==NULL)    // Empty tree? Make root!
    {
        Y = new node;
        Y->data=X;
        Y->LCH =Y->RCH=NULL;
        Y->BF   =0;
        root = Y;
        return; }
}
```

# The AVL Insert Process (In C): 2/8

```
// Locate insertion point for X.
// P scans through the tree
// Q is P's parent (Q lags behind P)
//     New Node Y will be a child of Q
// A is the last parent above Y w/BF=+/-1
// F is A's parent (F lags behind A)
//
F = Q = NULL;
A = P = root;

while (P != NULL) {
    if (P->BF != 0) {A=P; F=Q;}
    if (X == P->data) return;
    if (X < P->data) {Q=P; P=P->LCH;}
                    else {Q=P; P=P->RCH;}
}
```

# The AVL Insert Process (In C): 3/8

```
// At this point, P is NULL, but Q points
// at the last node where X belongs
// (either as Q's LCH or RCH)
//
Y = new node;
Y->data = X;
Y->LCH  = NULL; // New nodes are always
Y->RCH  = NULL; // inserted as leaves
Y->BF   = 0;     // Leaves always balanced

// Will Y be Q's new left or right child?
if (X < Q->data) Q->LCH = Y;
                else Q->RCH = Y;
```

# The AVL Insert Process (In C): 4/8

```
// Adjust BF's from A to Q. Since A was the
// LAST ancestor with a BF of +/- 1, ALL
// nodes BETWEEN A and Q must have a BF of
// 0, and will, therefore, BECOME +/-1.
// If X is inserted in the LEFT subtree
// of A, then d=+1 (d=-1 means we inserted
// X in the RIGHT subtree of A.
```

```
if (X > A->data) {P=A->RCH; B=P; d=-1;}
                else {P=A->LCH; B=P; d=+1;}
```

```
while (P != Y) {
    if (X>P->data) {P->BF=-1; P=P->RCH;}
    else {P->BF=+1; P=P->LCH;}
}
```

# The AVL Insert Process (In C): 5/8

```
// Now we check the BF at A and see if we  
// just BALANCED the tree, IMBALANCED the  
// tree, or if it is still BALANCED ENOUGH.
```

```
if (A->BF==0) {A->BF=d; // Tree is still  
                return;} // mostly balanced
```

```
if (A->BF+d==0) {A->BF=0; // Insertion put  
                return;} // tree INTO bal
```

```
// If we didn't take either of the two  
// returns immediately above, then the  
// tree is IMBALANCED. We have to  
// determine the required rotation type
```

# The AVL Insert Process (In C): 6/8

```
if (d==+1) { // left imbalance.  LL or LR?
    if (B->BF==+1) { // LL rotation
        // Change the child pointers at A and B
        // To reflect the rotation.  Adjust
        // the balance factors at A and B
        // <<< LEFT FOR YOU TO WRITE >>>
    }
    else { // LR Rotation
        // Adjust the child pointers of nodes
        // A, B, and C to reflect the new
        // structure after the rotation
        // <<< LEFT FOR YOU TO WRITE, >>>
        // <<< BUT HERE'S A HEAD START >>>
        C = B->RCH; // C is B's right child
        CL = C->LCH; // CL and CR are C's left
        CR = C->RCH; // and right children
    }
}
```



# The AVL Insert Process (In C): 7/8

```
switch (C->BF) {
    // Set the new balance factors at
    // A and B based on the BF at C
    // Note: There are 3 cases
    // <<< LEFT FOR YOU TO WRITE >>>
}

C->BF=0; B=C;
} end of else (LR Rotation)
} // end of "if (d=+1)"
else { // d=-1. This is a right imbalance
    // (RR or RL). THAT code goes here.
    // <<< LEFT FOR YOU TO WRITE >>>
}
```

# The AVL Insert Process (In C): 8/8

```
// The subtree with root B has been  
// reblanaced, and is the new subtree  
// of F. The original subtree of F had  
// root A.
```

```
// did we rebalance the root?  
if (F == NULL) {root=B; return;}
```

```
// otherwise, we rebalanced whatever was  
// the child (left or right) of F.  
if (A == F->LCH) {F->LCH=B; return;}  
if (A == F->RCH) {F->RCH=B; return;}  
cout << "We should never be here\n";
```

```
} // End of AVL_Insert
```

# Properties of AVL Trees

- Time to insert:  $O(h)$
- Time to search:  $O(h)$
- Time to find successor / predecessor:  $O(h)$
- Time to delete  $O(h)$
- Time to find min / max:  $O(h)$
- So, what is  $h$ ?
- AVL guarantees  $h \leq \sim 1.44 \lg n$ 
  - Where  $n$  is the number of nodes in the tree.
  - The 1.44 comes from Fibonacci theory

? Questions ?

---

? Questions ?