

CSI 403

ALGORITHMS AND DATA STRUCTURES

Lecture 14 – More Height-Balanced Trees:
Red-Black Trees (Deleting)

Review

- In BSTs (Binary Search Trees), the dynamic set operators (INSERT, DELETE, SEARCH, PREDECESSOR, SUCCESSOR, MAX, and MIN) all (can) take $O(h)$ time, where h is the height of the tree.
- In a height balanced tree, $h \approx \lg N$
- In a degenerate tree, $h = N$
- We need to have a way to keep the tree balanced as we perform repeated insertions and deletions.

Review, cont'd

◎ RBTs (Red-Black Trees):

- A different way of keeping a tree height-balanced
- All leaf nodes AND the root's parent are *T.nil*
- Five rules:
 - (1) All nodes are colored either “red” or “black”
 - (2) The root is black
 - (3) All leaves are black
 - (4) If a node is red, then both of its children are black
 - (5) For any node, the number of black nodes along any path to the leaves is the same as for any *other* such path to the leaves

Review, cont'd

- ◎ RBTs (Red-Black Trees):
 - BST property is always preserved
 - BST most of the operators work:
 - SEARCH, SUCCESSOR, PREDECESSOR, MIN, MAX
 - INSERT and DELETE are different than in “plain” BST
 - Fixing the tree requires *rotations*
 - RBT's have three different cases (kinds of problems) we may have to fix as the result of an insertion

Review, cont'd

⦿ Rotations

- Change the tree structure, but preserve the BST property
- Works by changing a limited (constant) number of pointers
 - “Local” adjustment
 - The amount of work done in a rotation is bound by some constant; *not* the height of the tree nor the number of nodes in it.

Review, cont'd

◎ RBT Insertion

- Do a “plain” search and insert the new node as red
- May violate one or more rules
- Do a fixup
 - May be in one of three cases
 - Fixing one case may solve the problem; may move us to another case ($1 \rightarrow 2 \rightarrow 3$).
 - Once we’ve fixed the problem(s) at any one level, move up two more levels and try again until we hit a level with no problems to fix. From there up to the root, the tree is valid.

Review, cont'd

◎ RBT Insertion – The Three cases:

- 1: z is red, $z.p$ is red, & z 's uncle y is red
 - Re-color z , z 's parent, and z 's uncle (no rotation)
- 2: z 's uncle y is black, z is a right child
 - Do a left rotation around z . Takes us immediately to
- 3: z 's uncle y is black, z is a left child
 - Recolor z 's parent and grandparent
 - Do a right rotation around z 's grandparent
- Move z up two levels ($z \leftarrow z$'s parent's parent), and try again as long as z 's parent is red

Deleting a Node from an RBT

⦿ RBT Deletion:

- Only slightly more complicated than insertion
- Similar to insert (at least from a high level)
 - Find node to delete
 - Delete / splice out (just like regular BST)
 - Clean up whatever we might have messed up
 - Cleanup runs in $O(\lg n)$ time
 - Leaves the tree balanced (h stays $O(\lg n)$)

RB-DELETE: the Pseudocode

```
RB-DELETE (T, z)           z points at the node to delete; it's not a key.
    if z.left == T.nil or z.right == T.nil
        y = z
    else y = TREE-PREDECESSOR(z)
    if y.left != T.nil then x = y.left else x = y.right
    x.p = y.p
    if y.p == T.nil
        T.root = x
    else if y == y.p.left
        y.p.left = x
    else y.p.right = x
    if y != z
        z.key = y.key //also copy y's satellite data (if any) into z
    if y.color == BLACK then RB-DELETE-FIXUP(T, x)
Return(y)
```

RB-DELETE-FIXUP: the Pseudocode

```
RB-DELETE-FIXUP (T, x)
while x ≠ T.root and x.color == BLACK
    if x == x.p.left
        w = x.p.right
        if w.color = RED
            w.color = BLACK
            x.p.color = RED
            LEFT-ROTATE(T, x.p)
            w = x.p.right
        if w.left.color == black and w.right.color == BLACK
            w.color = RED
            x = x.p
        else if w.right.color == BLACK
            w.left.color = BLACK
            w.color = RED
            RIGHT-ROTATE(T, w)
            w = x.p.right
            w.color = x.p.color
            x.p.color = BLACK
            w.right.color = BLACK
            LEFT-ROTATE(T, x.p)
            x = T.root
    else (same as then clause with "right" and "left" exchanged)
        x.color = BLACK
```

CASE 1

CASE 2

CASE 3

CASE 4

RB-DELETE-FIXUP: the Pseudocode

```
RB-DELETE-FIXUP( $T, x$ )
1  while  $x \neq T.root$  and  $x.color == BLACK$ 
2      if  $x == x.p.left$            // is  $x$  a left child?
3           $w = x.p.right$          //  $w$  is  $x$ 's sibling
4          if  $w.color == RED$ 
5               $w.color = BLACK$ 
6               $x.p.color = RED$ 
7              LEFT-ROTATE( $T, x.p$ )
8               $w = x.p.right$ 
9          if  $w.left.color == BLACK$  and  $w.right.color == BLACK$ 
10              $w.color = RED$ 
11              $x = x.p$ 
12         else
13             if  $w.right.color == BLACK$ 
14                  $w.left.color = BLACK$ 
15                  $w.color = RED$ 
16                 RIGHT-ROTATE( $T, w$ )
17                  $w = x.p.right$ 
18              $w.color = x.p.color$ 
19              $x.p.color = BLACK$ 
20              $w.right.color = BLACK$ 
21             LEFT-ROTATE( $T, x.p$ )
22              $x = T.root$ 
23     else // same as lines 3–22, but with “right” and “left” exchanged
24          $w = x.p.left$ 
25         if  $w.color == RED$ 
26              $w.color = BLACK$ 
27              $x.p.color = RED$ 
28             RIGHT-ROTATE( $T, x.p$ )
29              $w = x.p.left$ 
30         if  $w.right.color == BLACK$  and  $w.left.color == BLACK$ 
31              $w.color = RED$ 
32              $x = x.p$ 
33         else
34             if  $w.left.color == BLACK$ 
35                  $w.right.color = BLACK$ 
36                  $w.color = RED$ 
37                 LEFT-ROTATE( $T, w$ )
38                  $w = x.p.left$ 
39              $w.color = x.p.color$ 
40              $x.p.color = BLACK$ 
41              $w.left.color = BLACK$ 
42             RIGHT-ROTATE( $T, x.p$ )
43              $x = T.root$ 
44      $x.color = BLACK$ 
```

Deleting a Node from an RBT

- In $\text{RB-DELETE-FIXUP}(T, x)$, what is x ?
 - “The node passed to RB-DELETE-FIXUP is one of two nodes: either the node that was y ’s sole child before y [from RB-DELETE] was spliced out if y had a child that was not the sentinel $T.nil$, or, if y had no children, x is the sentinel $T.nil$. In the latter case, the unconditional assignment in line 7 ($x.p = y.p$) guarantees that x ’s parent is now the node that was previously y ’s parent, whether x is a key-bearing internal node, or the sentinel, $T.nil$.”
 - In other words,
 - y is the node that was spliced out
 - x is either:
 - y ’s sole non-sentinel child before y was spliced out, or
 - The sentinel if y had no children.

Deleting a Node from an RBT

⦿ If y is black, what could we have violated?

1) All nodes are colored either “red” or “black”

2) The root is black

If y is the root and x is red, then the root has become red.

3) All leaves are black

4) If a node is red, then its children are both black
if x and $y.p$ are both red, then #4 is violated.

Deleting a Node from an RBT

- If y is black, what could we have violated?
 - 5) For any node, the black height is the same for all paths to the leaves

Any path containing y now has 1 fewer black nodes.

Fix by giving x an “extra black”

Add 1 to the black height on paths containing x

Now #5 is OK, but #1 is not.

x is now either *doubly black* (if $x.\text{color} == \text{BLACK}$) or *red and black* (if $x.\text{color} == \text{RED}$)

RB-DELETE-FIXUP

- ⦿ The idea behind the fixup routine:
- ⦿ Move the extra black up the tree until
 - x points to a red & black node \rightarrow turn it into a black node,
 - x points to the root \rightarrow just remove the extra black, or
 - we can do certain rotations and re-colorings and finish.

Within the `while` loop:

x always points to a non-root doubly-black node.


w is x 's sibling.

w cannot be $T.nil$, since that would violate property 5 at $x.p$.

RB-DELETE-FIXUP (2)


- There are 8 cases, 4 of which are symmetric to the other 4.
- As with insertion, the cases are not mutually exclusive (fixing one case may create a problem that falls into another case).
- We'll look at cases in which x is a left child.

RB-DELETE-FIXUP (3)



Case 1

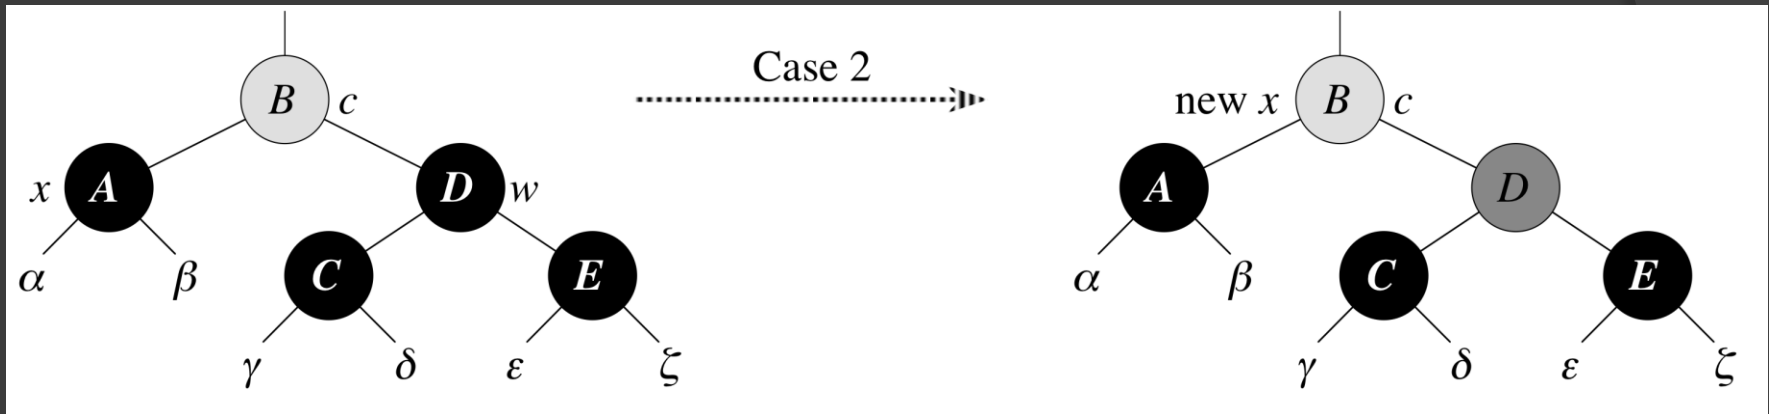
- w is red and must have black children (otherwise property 4 violated)



- w must have black children
- Make w black and $x.p$ red
- Then left rotate on $x.p$
- New sibling of x was a child of w before rotation, so it must be black
- Go immediately to case 2, 3, or 4

Case 2


- w is black and both of w 's children are black



- Take 1 black off x (singly black) and off w (red)
- Move that black to $x.p$
- Do the next iteration with $x.p$ as the new x
- If we entered this case from case 1, then $x.p$ was red
- New x is red & black \rightarrow color of new x is RED \rightarrow loop terminates. The new x is made black in the last line

Case 3


- w is black, w 's left child is red, and w 's right child is black



- Make w red and w 's left child black
- Then right rotate on w
- New sibling w of x is black with a red right child. Go on immediately to Case 4

Case 4

- w is black, w 's left child is black, and w 's right child is red



- Make w be $x.p$'s color (c)
- Make $x.p$ black and w 's right child black
- Then left rotate on $x.p$
- Remove the extra black on x (x is now singly black) without violating any red-black properties
- All done. Setting x to root causes the loop to terminate

Analysis of Deletion

- ⦿ $O(\lg n)$ time to get through RB-DELETE up to the call of RB-DELETE-FIXUP
- ⦿ Within RB-DELETE-FIXUP:
 - Case 2 is the only case where more iterations occur
 - x moves up 1 level
 - Hence, no more than $O(\lg n)$ iterations
 - Cases 1, 3, & 4 have 1 rotation each, so there can't be any more than 3 rotations in all
- ⦿ Hence, $O(\lg n)$ time (total)